



**R
I
SE**

CONTINUOUS DEEP ANALYTICS (2018-2023)





ROYAL INSTITUTE
OF TECHNOLOGY

The Team

Lars
Kroll



Distributed
Computing
Systems

*prog.languages
distr. computing*

Paris
Carbone



Christian
Schulte



Seif
Haridi



New MSc/PhDs

- **Klas Segeljakt**
- **Max Meldrum**
- **Oscar Bjühr**
- **Johan Mickos**

Theodore
Vasiloudis



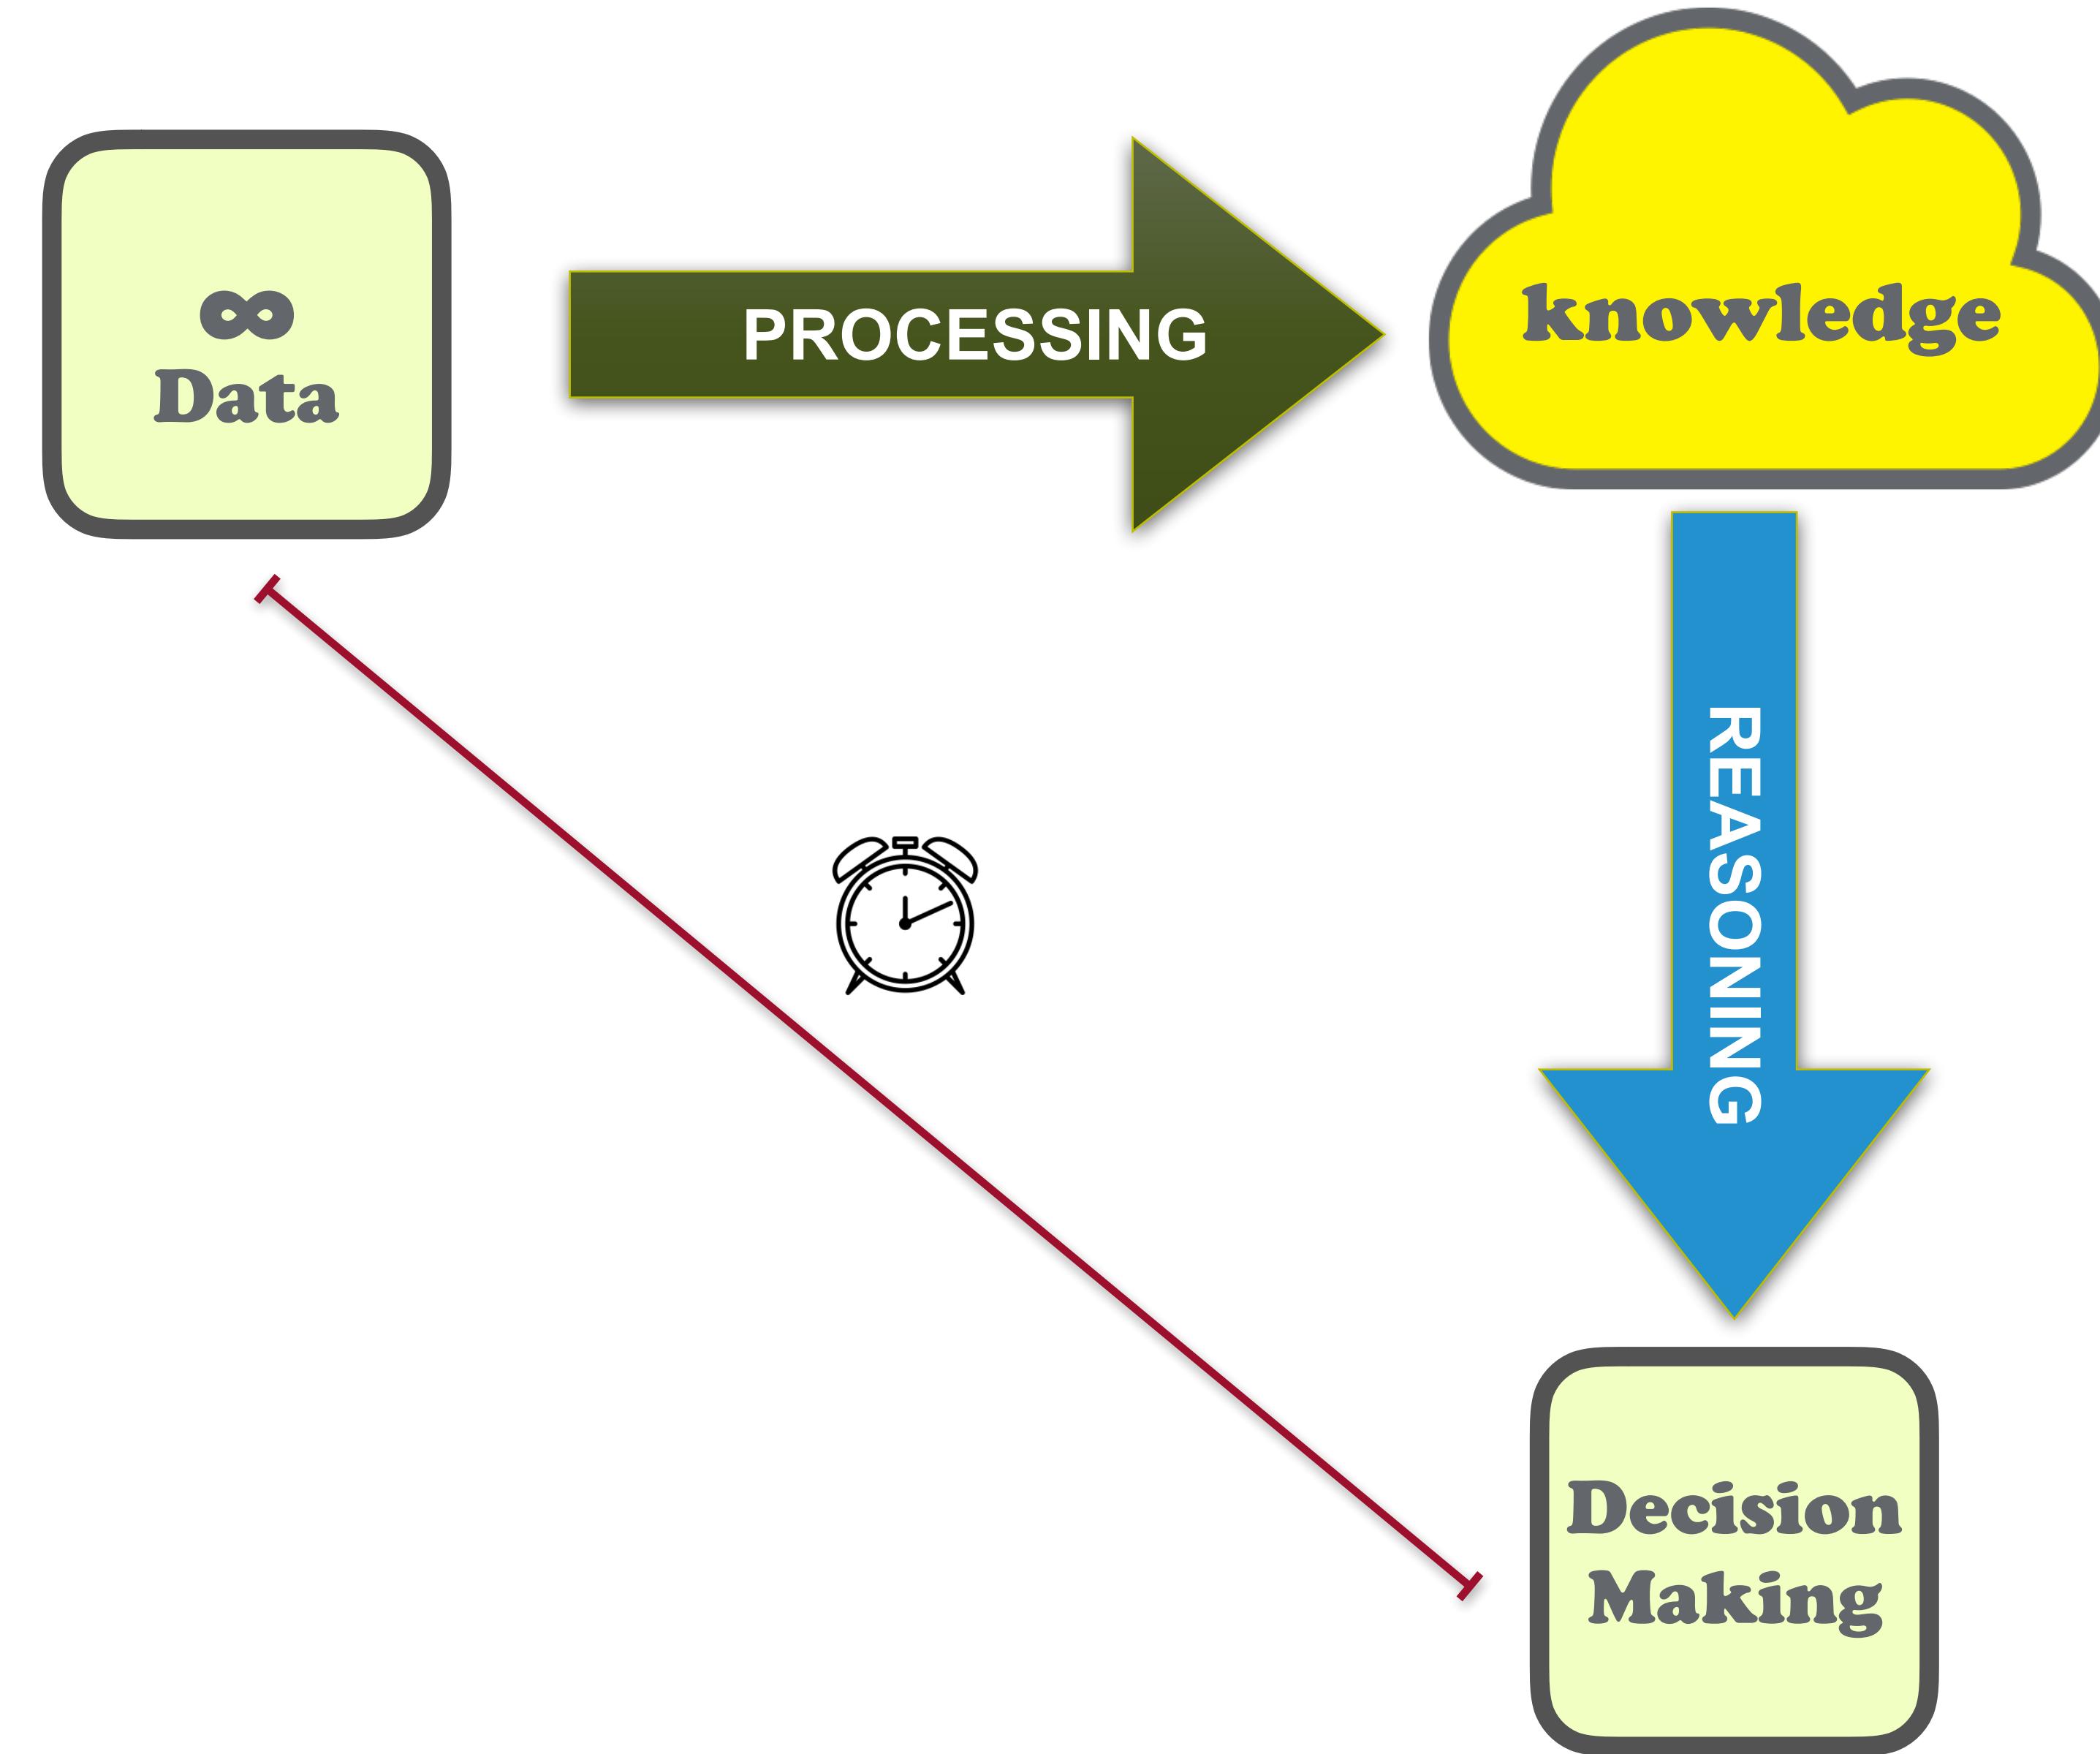
Machine
Learning
Expertise

Daniel
Gillblad





ROYAL INSTITUTE
OF TECHNOLOGY





ROYAL INSTITUTE OF TECHNOLOGY

The diagram illustrates a two-stage feature learning process. In the first stage, three input images (a noisy grid, a face sketch, and a grayscale face) are processed through three parallel layers of linear transformations (represented by green boxes with sigmoid activation functions). The outputs of these layers are then concatenated into a single vector. In the second stage, this concatenated vector is processed through four more layers of linear transformations with sigmoid activations. The final output is a probability value representing the likelihood that the voice is male.

inputs for acoustic properties

x_0

x_1

x_2

x_3

x_4

x_5

\vdots

x_{19}

w_0

w_1

w_2

w_3

w_4

w_5

w_{19}

weights

sum

$+$

b bias

sigmoid

probability voice is male

Linear Transformation

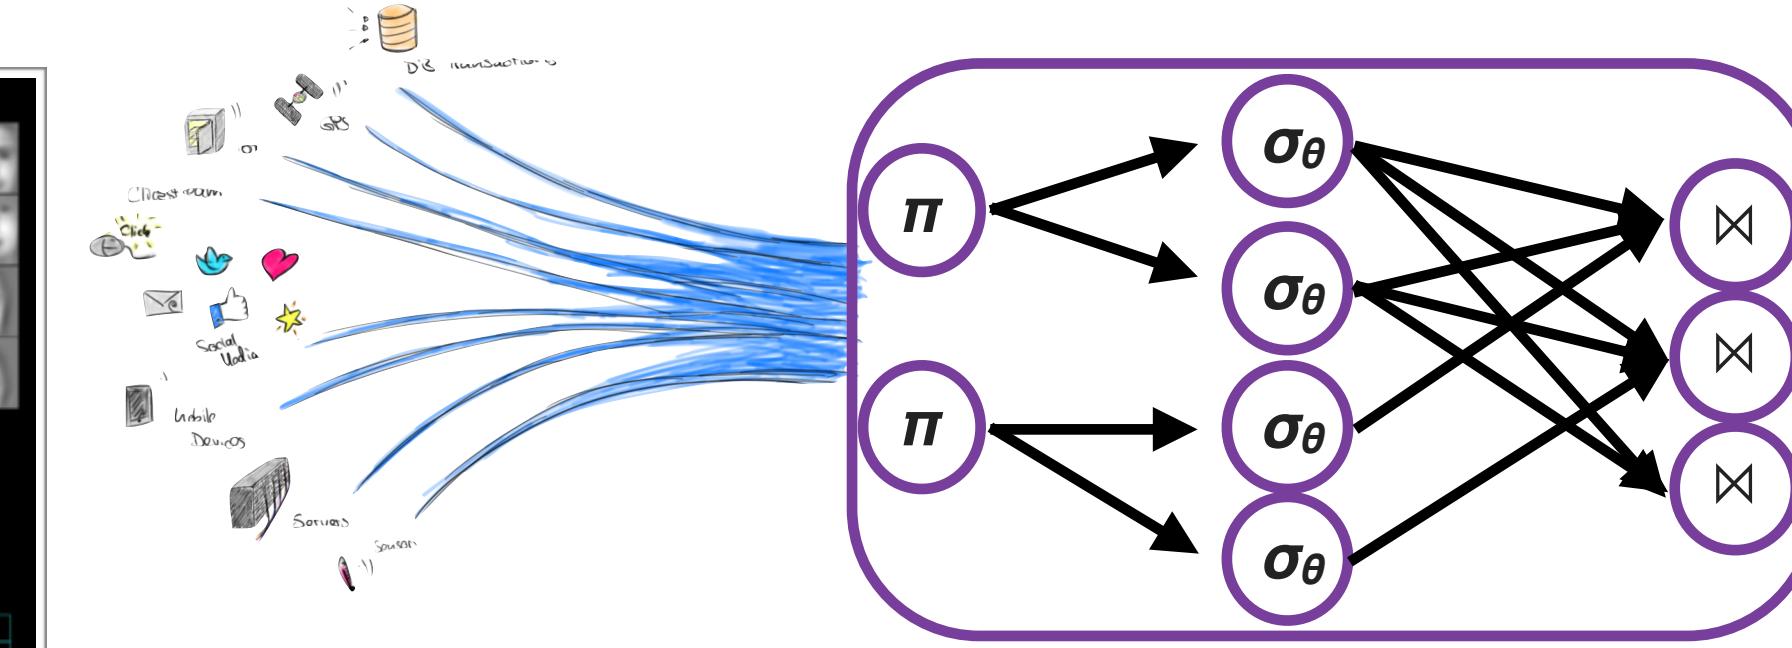
Linear Transformation

Linear Transformation

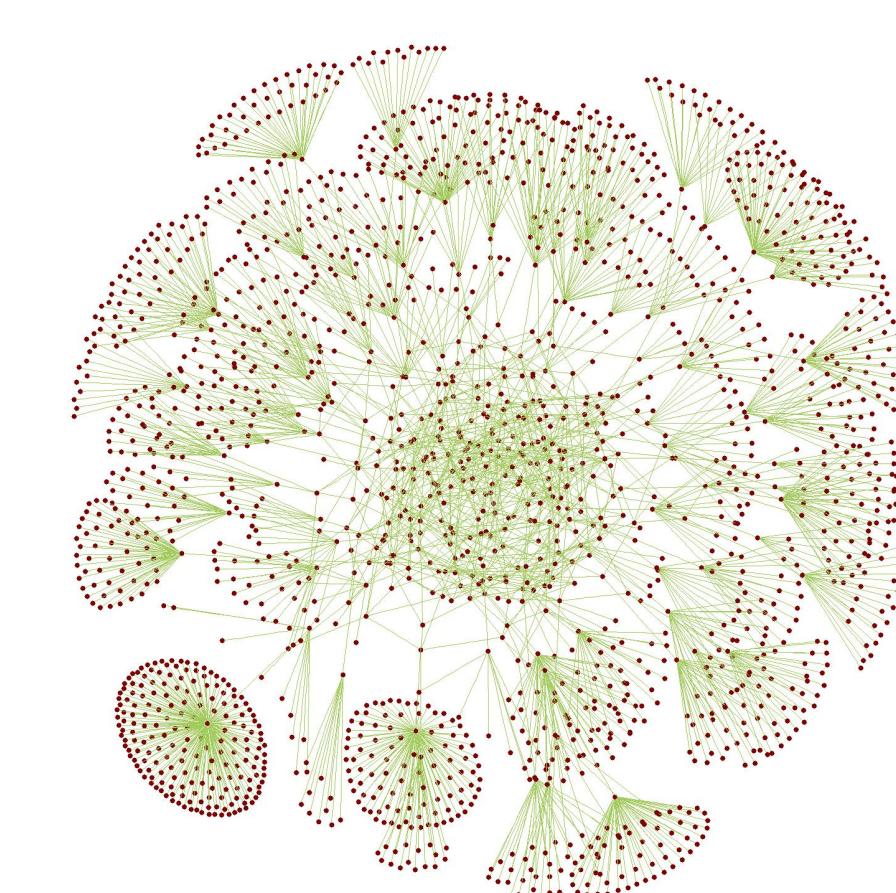
Linear Transformation

Feature Learning

Tensor Programming



Relational Data Streams

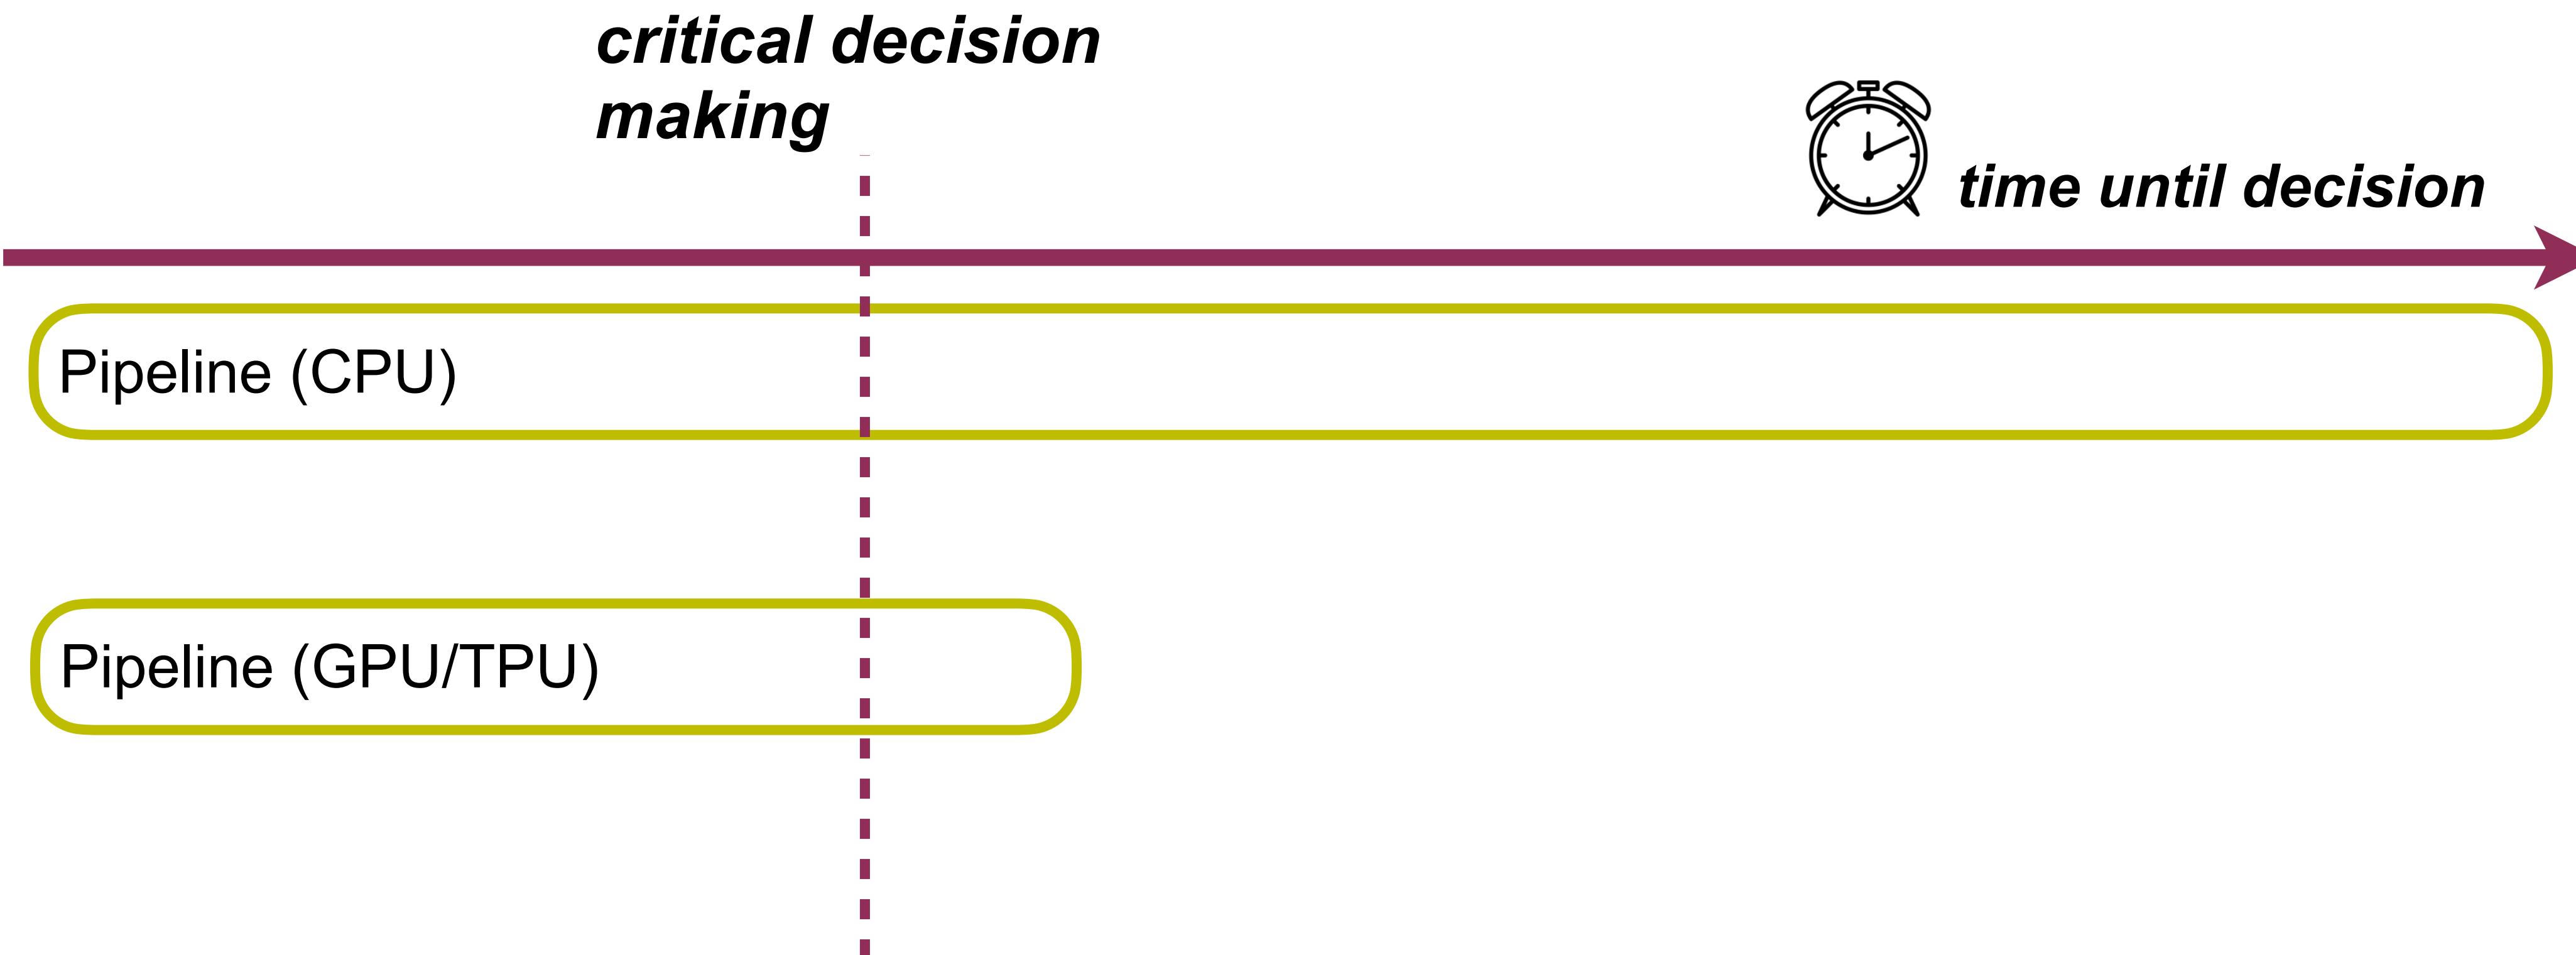


Dynamic Graphs



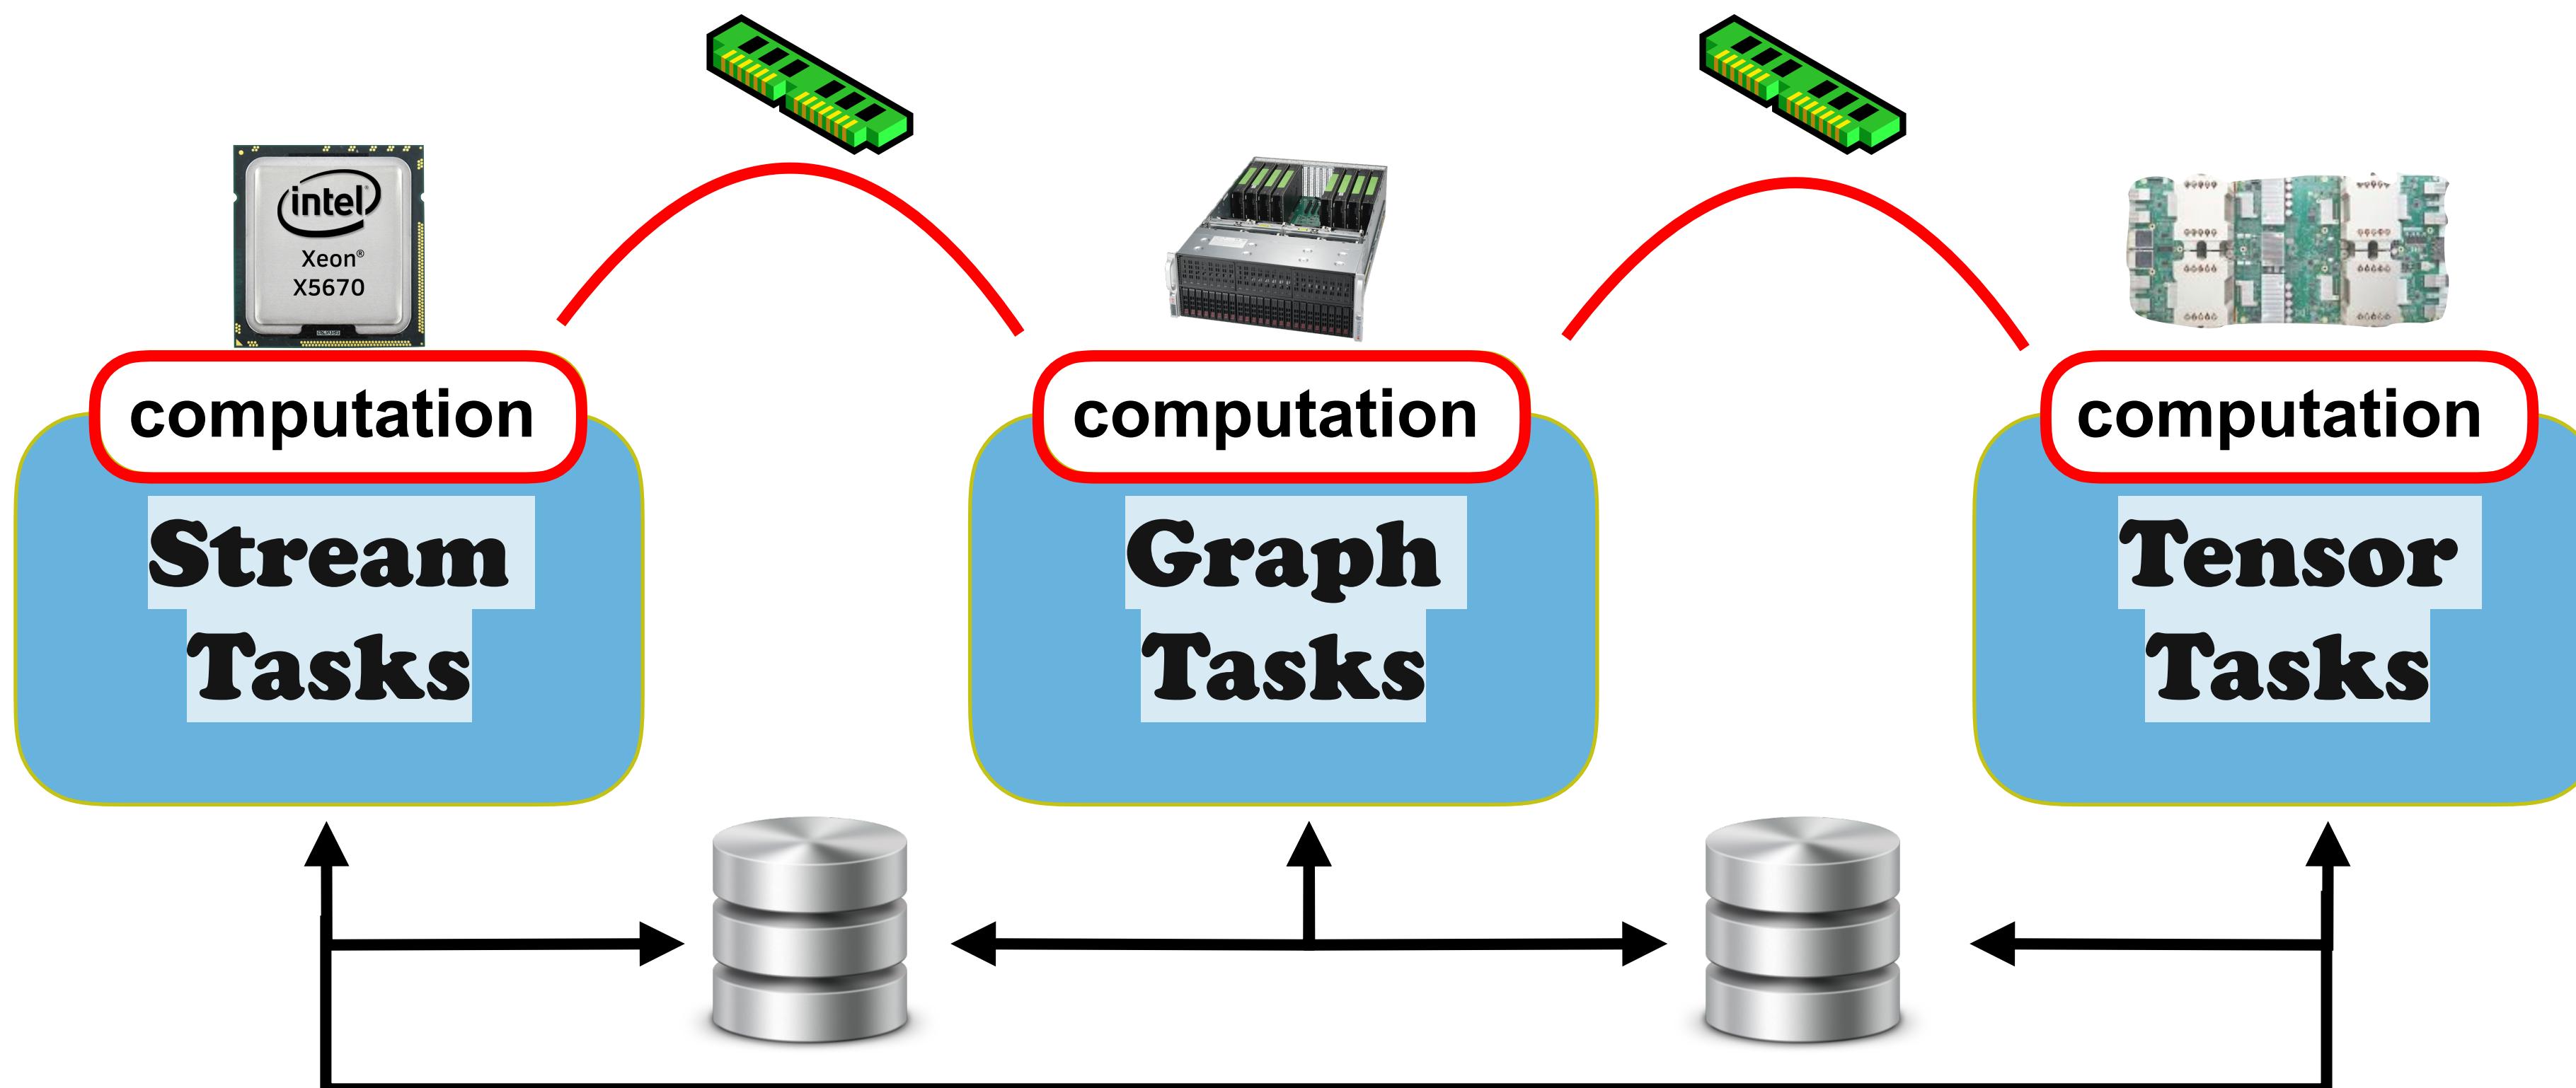
ROYAL INSTITUTE
OF TECHNOLOGY

Hardware Acceleration: Important but Not Enough



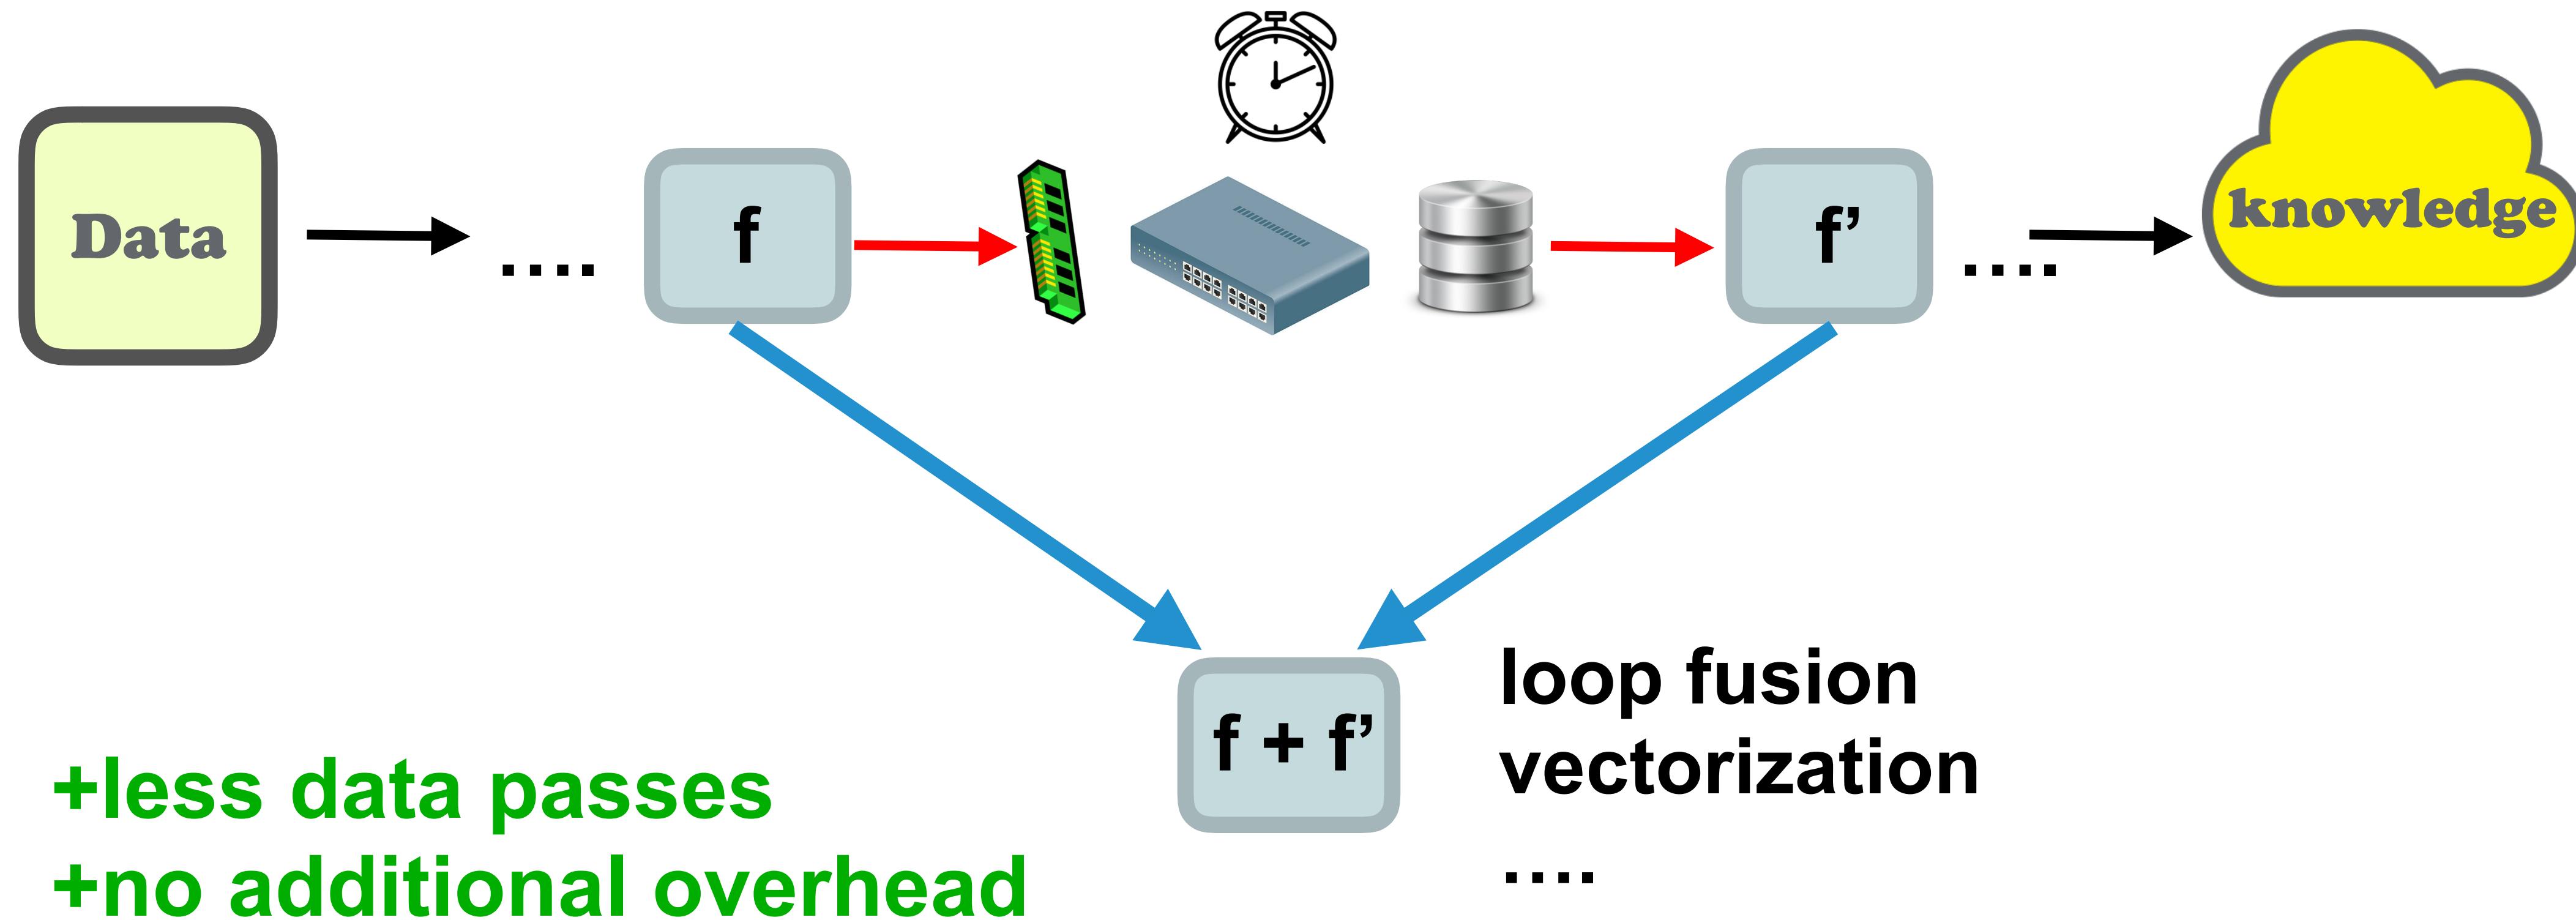
Cross-Platform Computation is Inefficient

- No computation sharing optimisations



- expensive data exchange through disk

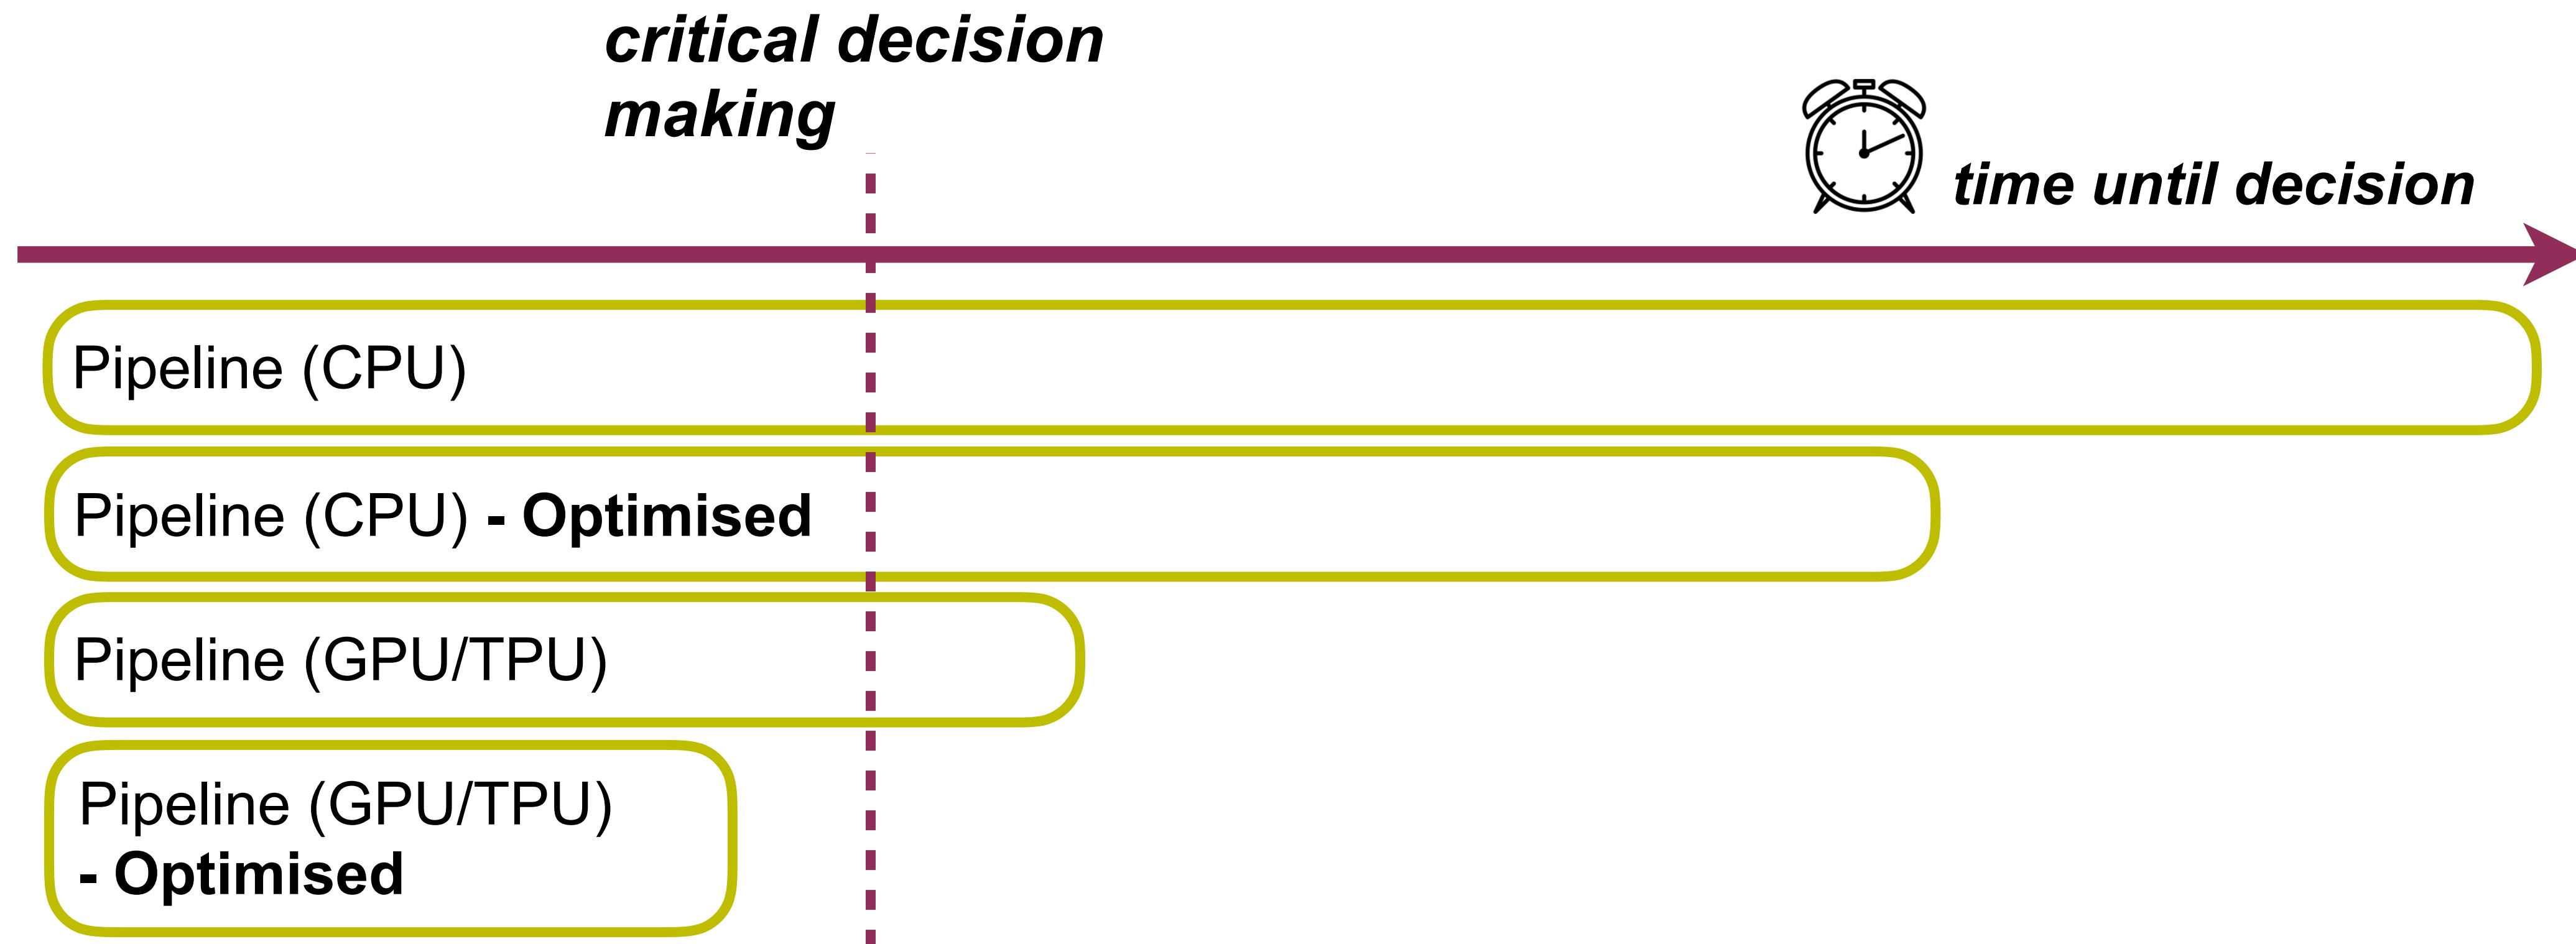
Computation Sharing





ROYAL INSTITUTE
OF TECHNOLOGY

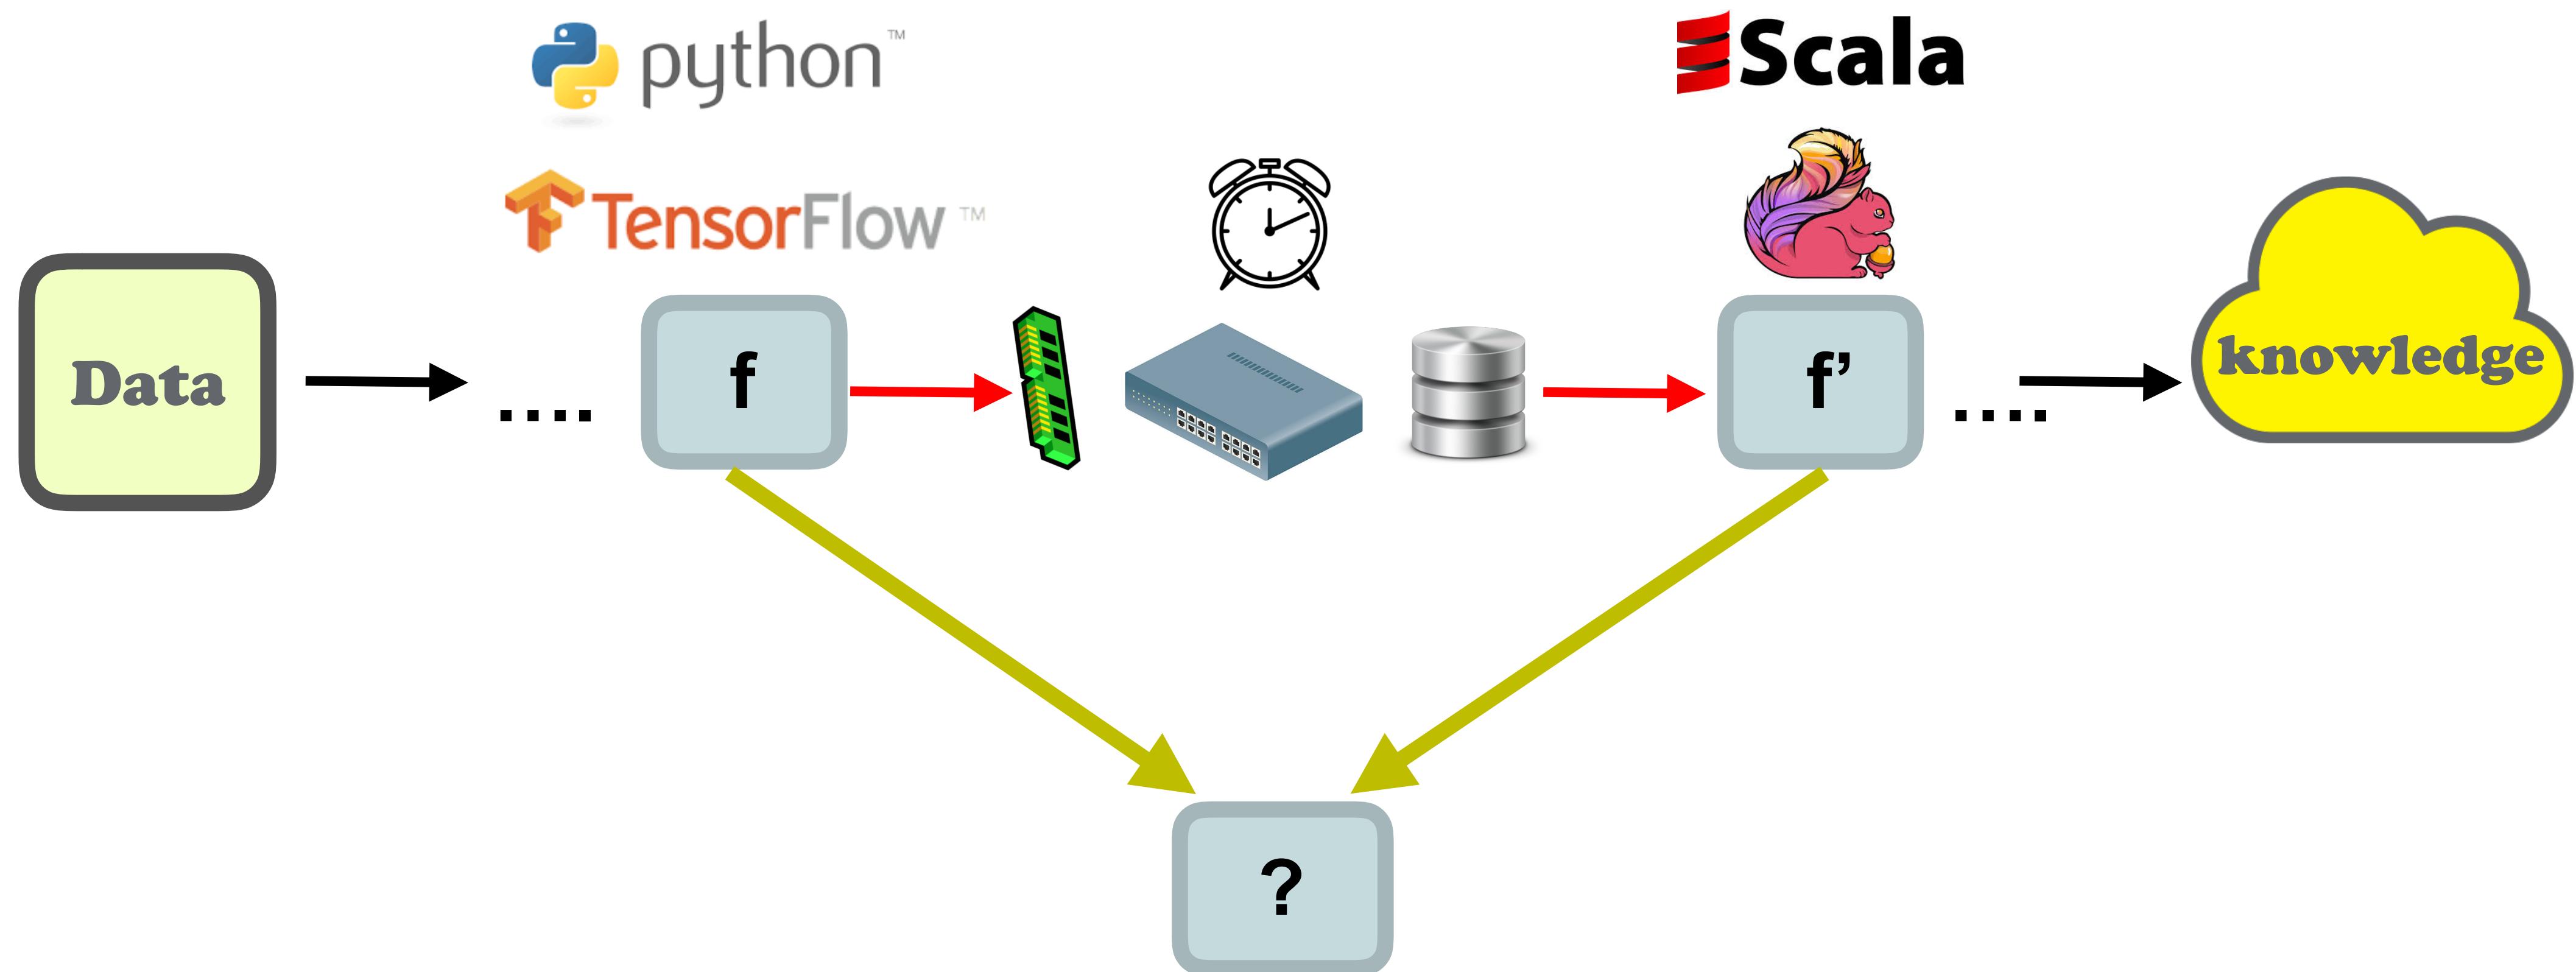
Critical Decision Making demands Efficiency





ROYAL INSTITUTE
OF TECHNOLOGY

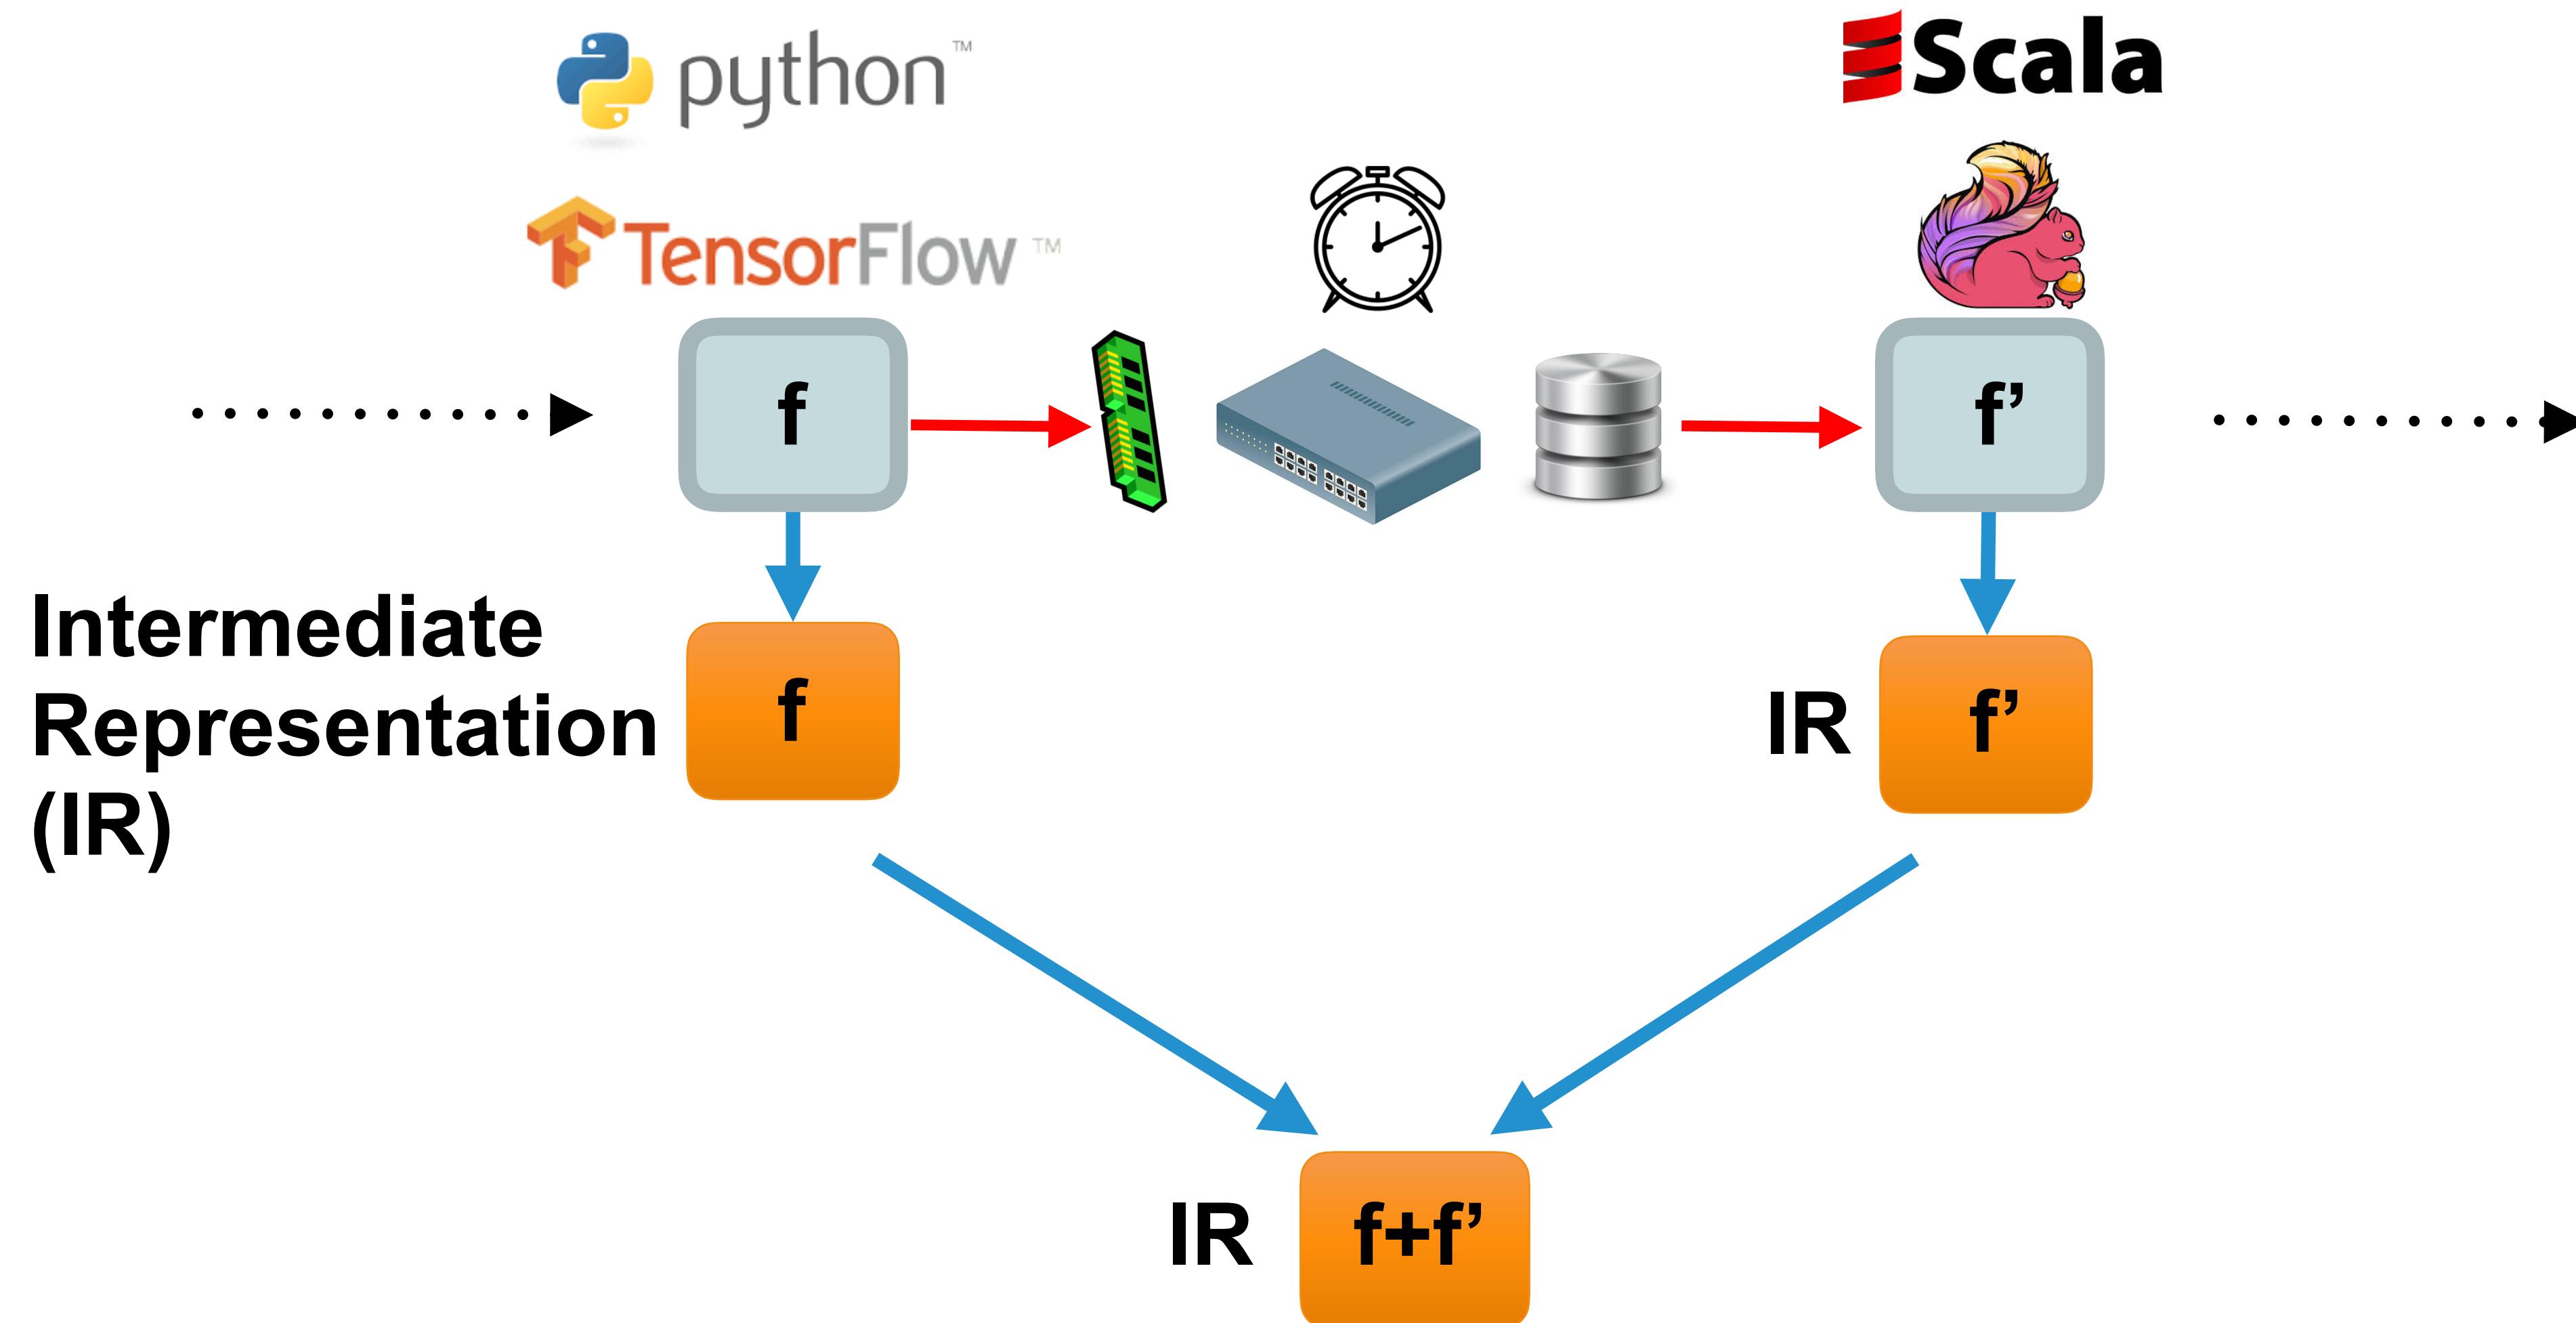
The Problem





ROYAL INSTITUTE
OF TECHNOLOGY

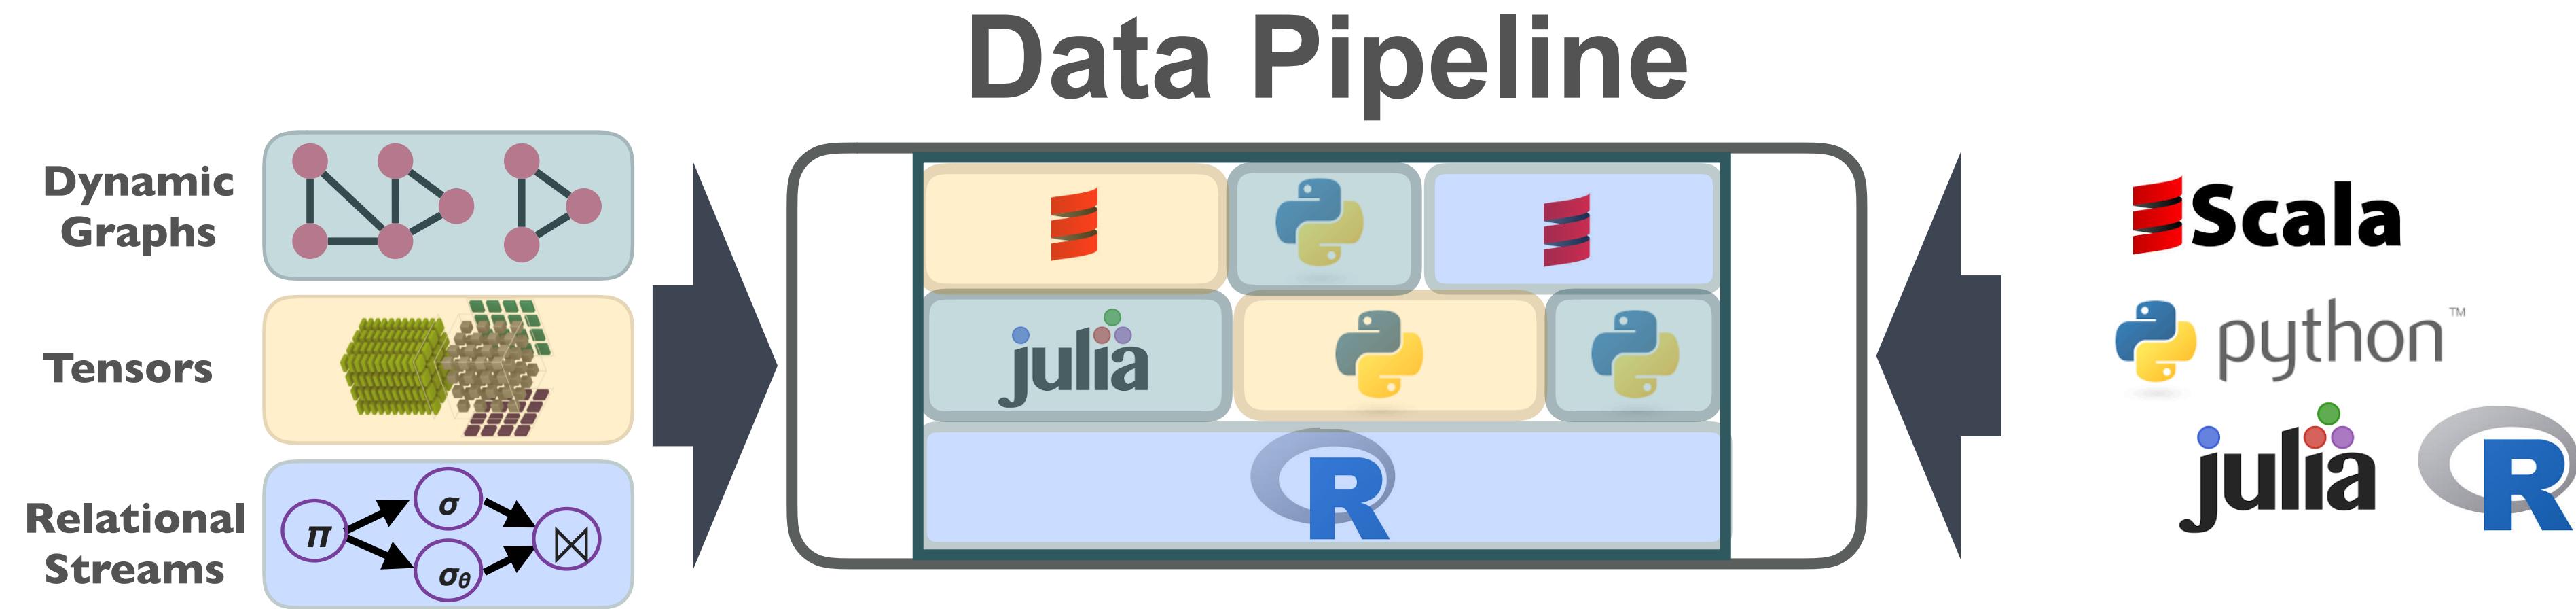
The Solution





ROYAL INSTITUTE
OF TECHNOLOGY

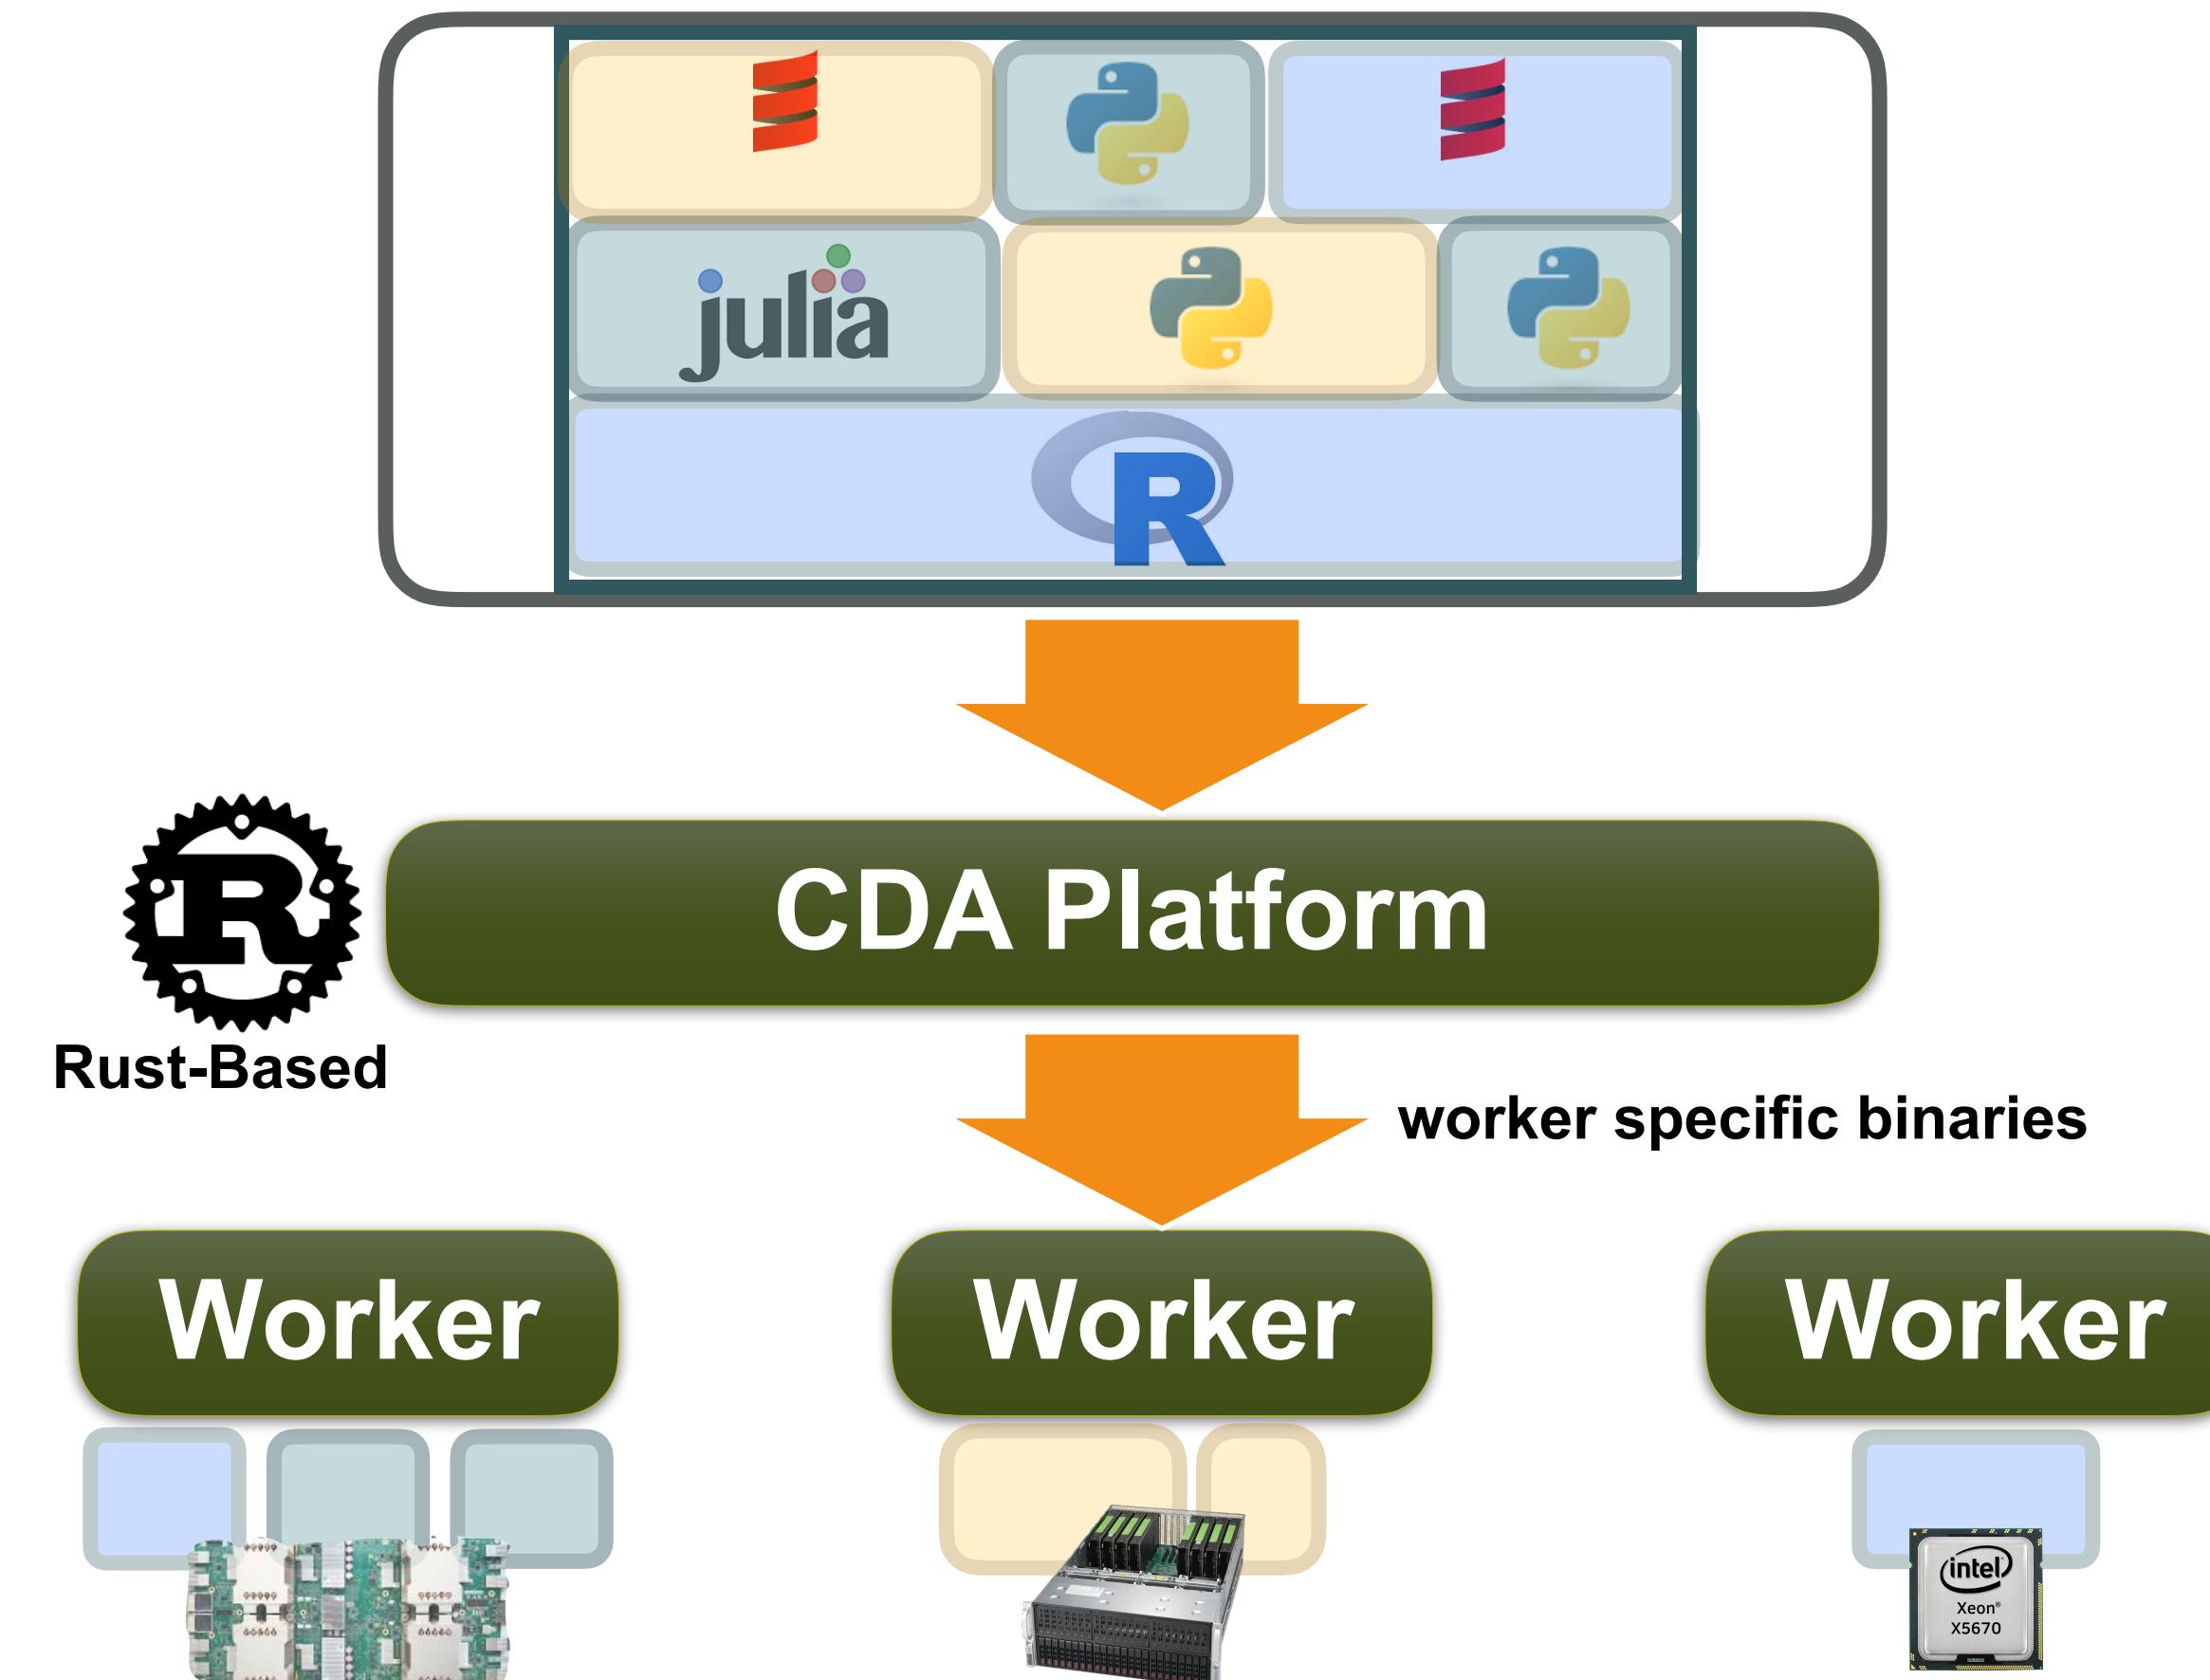
Model and Language Independence





ROYAL INSTITUTE
OF TECHNOLOGY

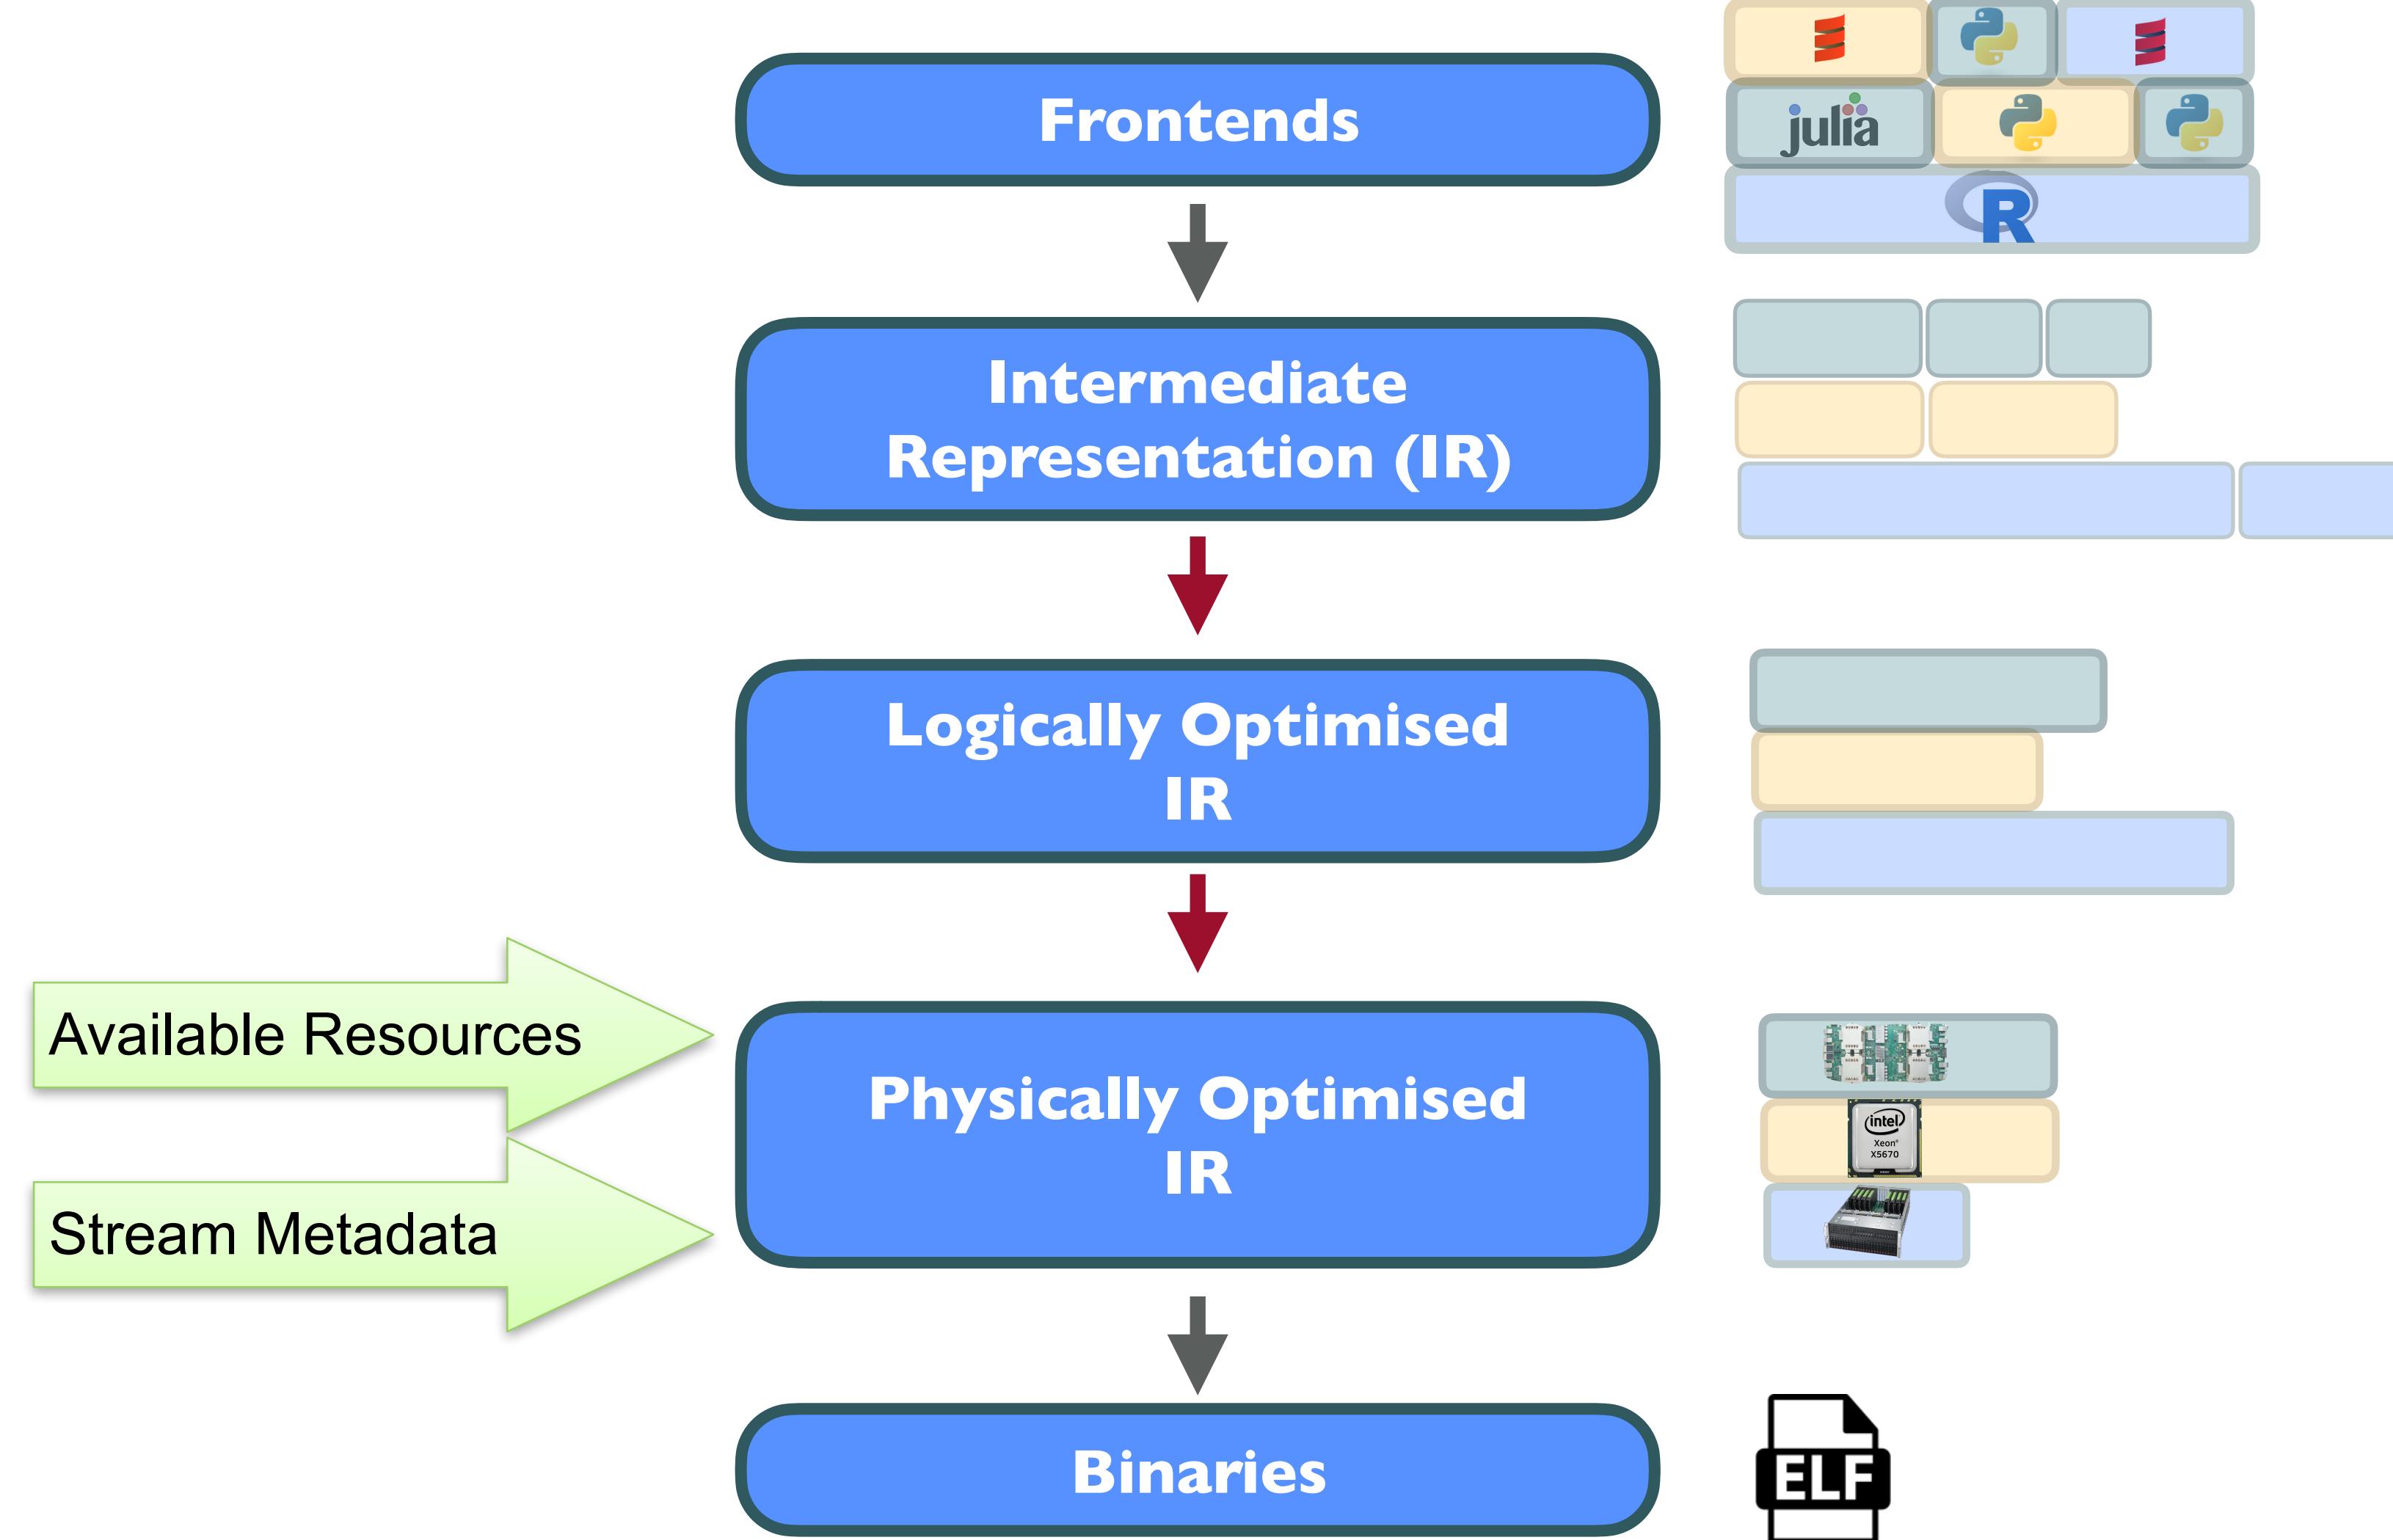
A Distributed Runtime for Heterogeneous HW





ROYAL INSTITUTE
OF TECHNOLOGY

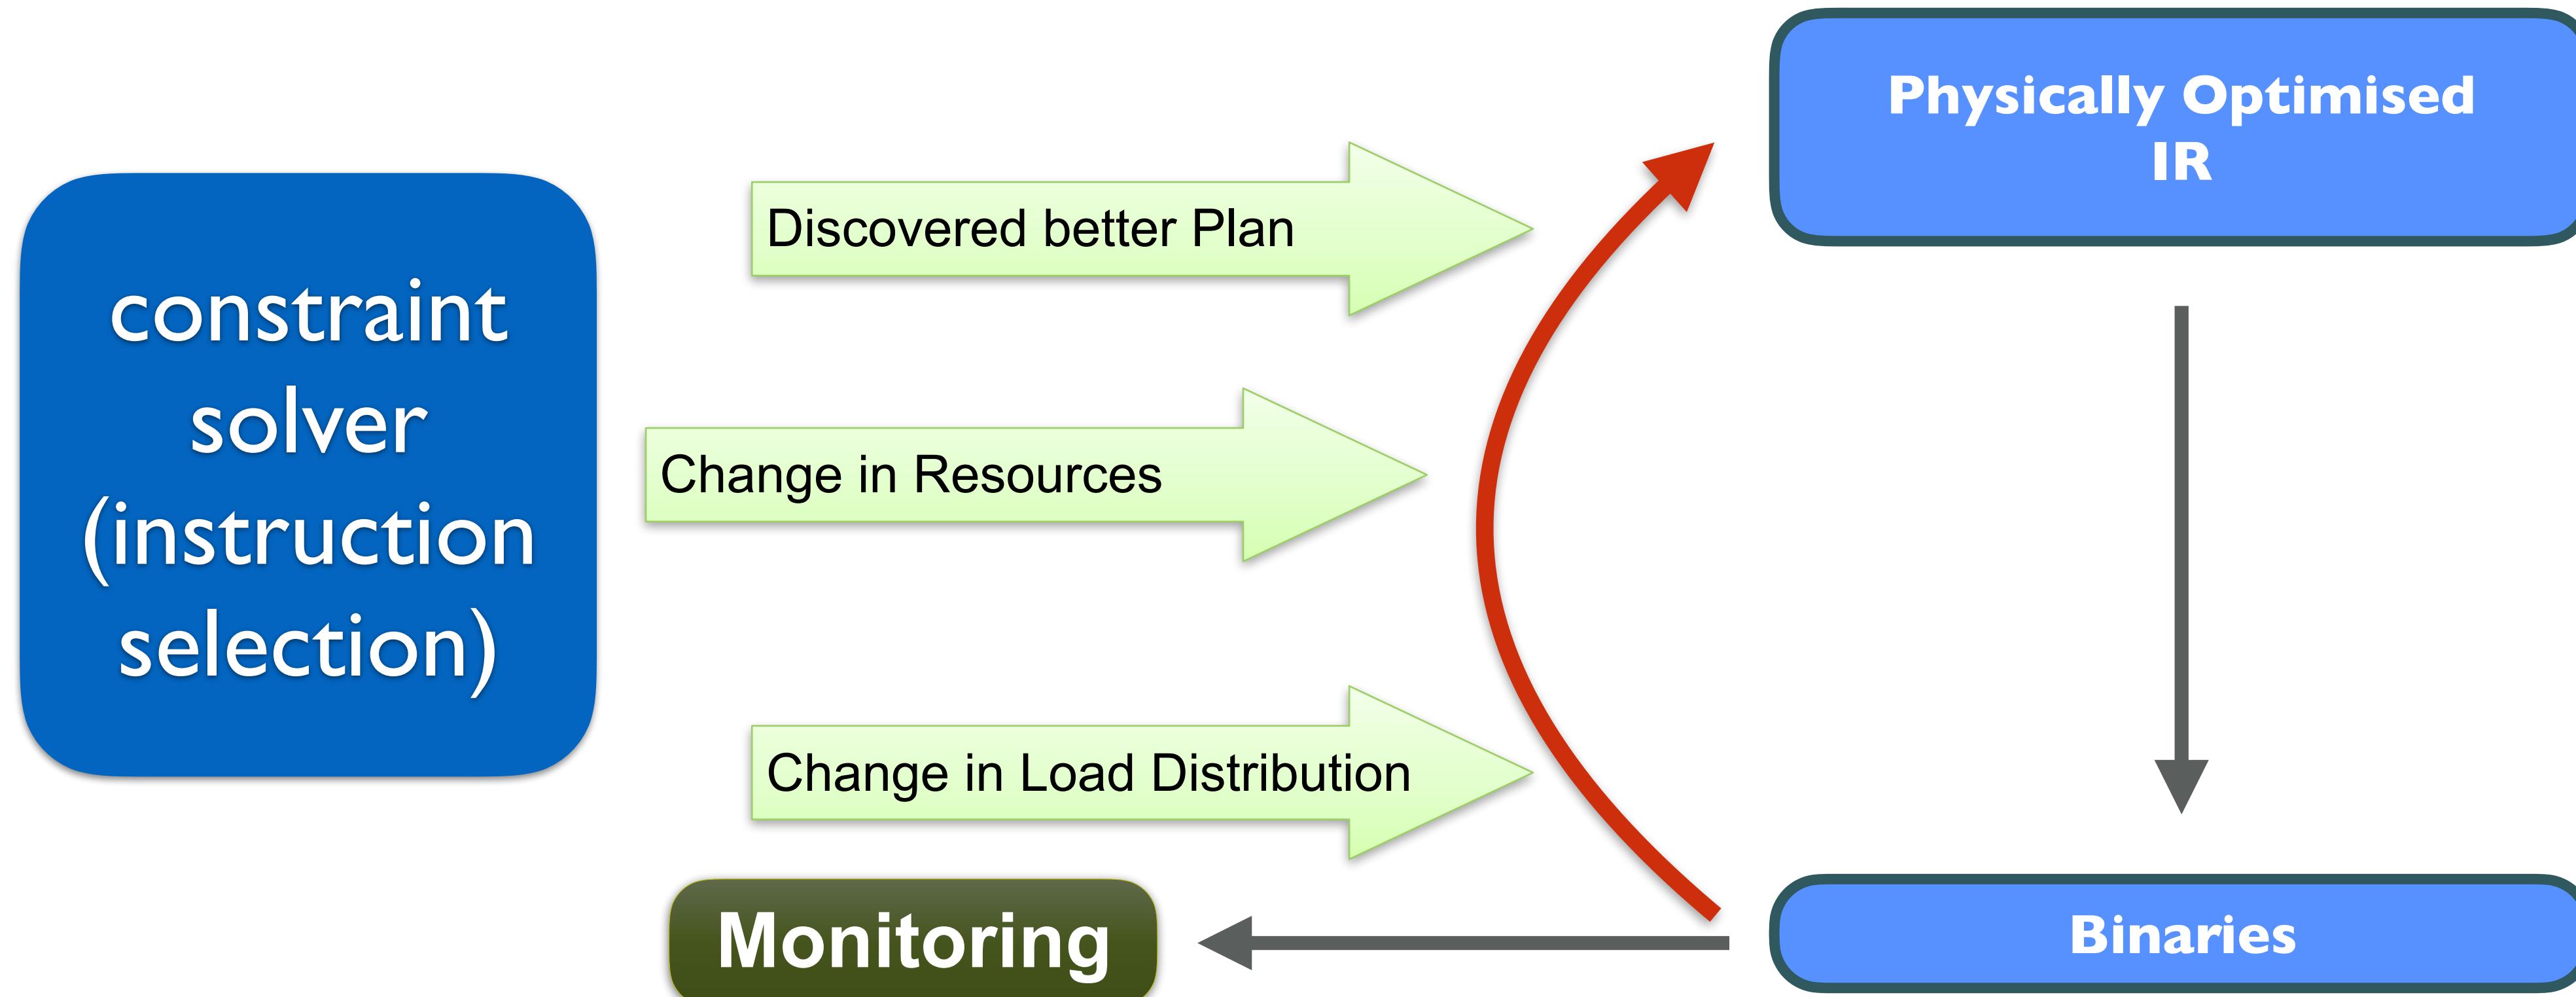
Target Architecture





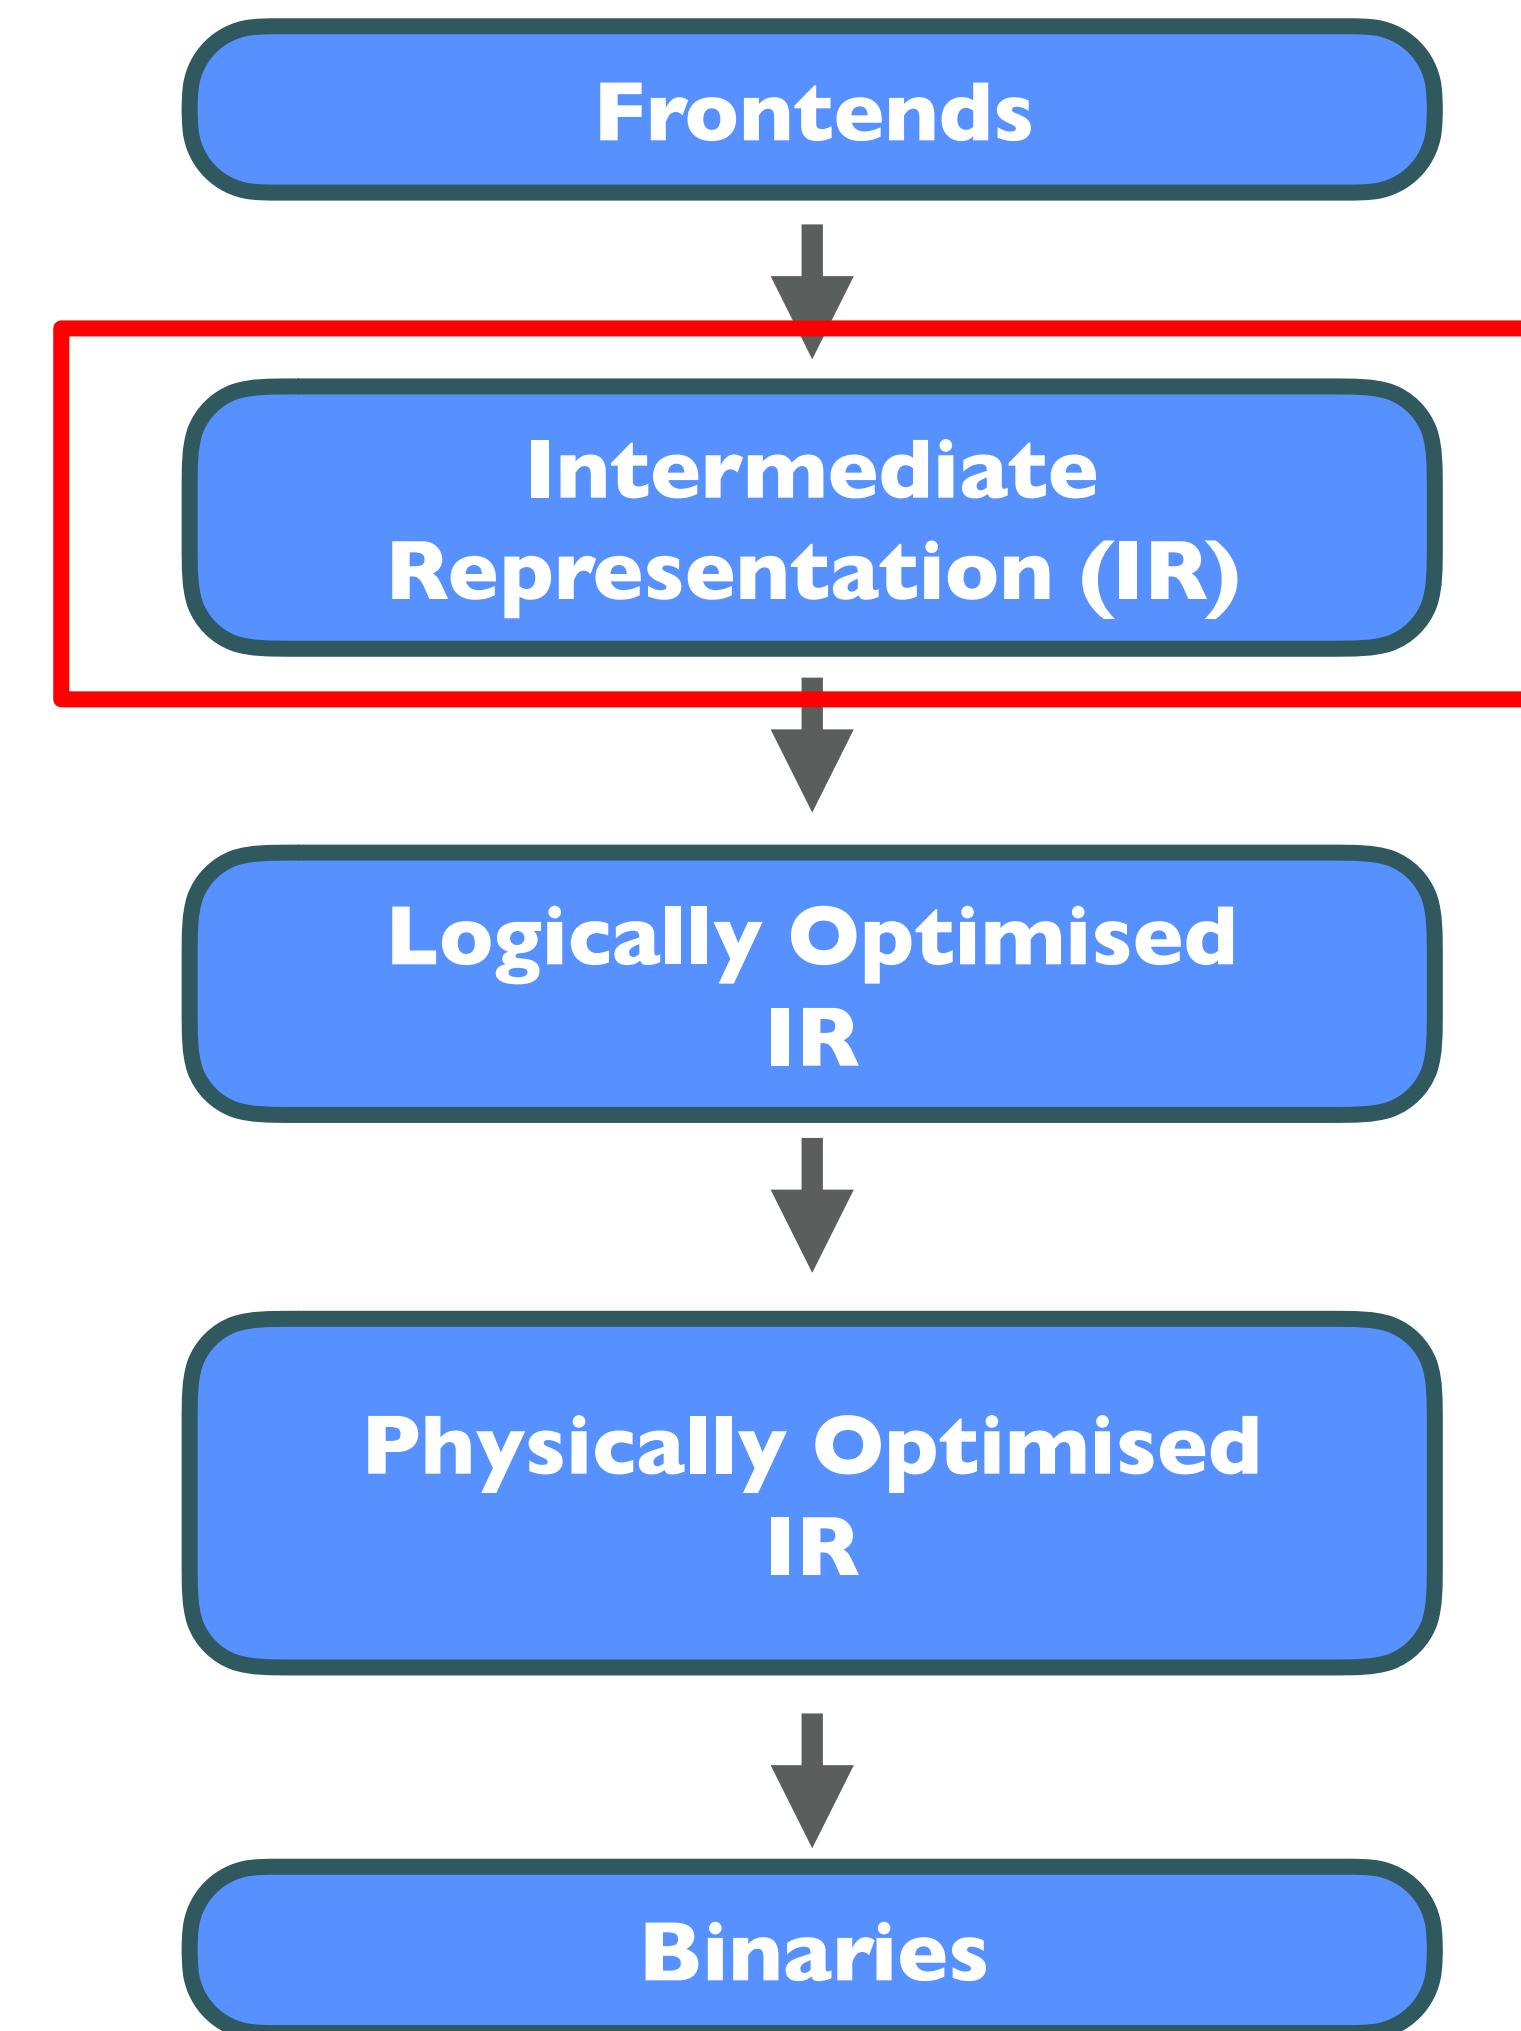
ROYAL INSTITUTE
OF TECHNOLOGY

Distributed JIT Compilation





ROYAL INSTITUTE
OF TECHNOLOGY





ROYAL INSTITUTE
OF TECHNOLOGY

From Weld to Arc

Extending the Weld IR for Streaming

Lars Kroll - KTH

What is Weld?



- A restrictive language for describing data transformations
- Pure expressions without side effects
- A compiler that produces LLVM IR
- Compiler leverages Weld's declarative syntax to make optimisations

Weld IR - Types

- Scalar+SIMD types
 - bool, u8...u64, i8...i64, f32, f64
- **Collections:** Read-only data types
 - vec, dict
- **Builders:** Write-only data types
 - appender, merger, groupbuilder
 - additive monads
- Structs { ... }: more like Tuples really
 - Builders are compositional over structs
 - I.e. structs of builders are also builders

No read-write
data types!

Weld IR - Ops

- `for` is a parallel loop over a collection (or iterator) and into a builder
- `merge` consumes a builder and a value and produces a new builder with the value merged in (according to the builder's semantics)
- `result` turns a builder into the corresponding type, i.e.
 - `appender[i32]` into `vec[i32]`
 - `merger[i32, +]` into `i32` (sum)
- In Weld you may only call `result` on a builder once and it consumes the builder (*linear type*)
- `if`, `lookup`, math functions, binary ops, casts, c-udfs, etc.

Weld Compilation

- Online (~JIT) compilation
- Quick and easy *type inference* to support dynamically typed front-ends (e.g., Python)
- Monadic properties allow automatic *parallelisation* and *vectorisation*
- Declarative style allows data access optimisations, such as *loop fusion* and *filter reordering*

Weld Example

Scala `input.map(i: Int => i + 5)`

Weld `| input:vec[i32] |
result(
 for(input,
 appender[i32],
 | app, _, i |
 merge(app, i + 5))))`

Weld Example

Scala `input.map(i: Int => i + 5)`

Weld `| input:vec[i32] |
result(
 for(input:vec[i32],
 appender[i32],
 | app:appender[i32], _:i64, i:i32 |
 merge(app, i + 5))))`



ROYAL INSTITUTE
OF TECHNOLOGY

Arc

- Arc extends Weld for streaming
- Observation
 - Stream Sources are **read-only**
 - Stream Sinks are **write-only**
 - Connect Sinks to Sources via *Channels*
- Source is a **collection** stream[T]
- Sink is a **builder** streamappender[T]
- Calling result on a Sink returns a Source and creates a Channel between them



ROYAL INSTITUTE
OF TECHNOLOGY

Arc Example I

Scala `input.map(i: Int => i + 5)`

Arc `| source: stream[i32], sink: streamappender[i32] |
for (source,
 sink,
 | out, i |
 merge(out, i + 5))`

Arc Example 2

Scala `input.filter(i: Int => i > 5)`

Arc `| source:stream[i32], sink: streamappender[i32] |
for(source,
 sink,
 |out, i|
 if (i > 5, merge(out, i), out))`

Arc Example 3

Scala

```
val mapped = input.map(i: Int => i + 5)
mapped.filter(i: Int => i % 2 == 0).toSink(...)
mapped.filter(i: Int => i % 2 != 0).toSink(...)
```

Arc

```
| source:stream[i32],
evenSink:streamappender[i32],
oddSink:streamappender[i32] |
let mapped = result(for(source,
    streamappender[i32],
    | out, i| merge(out, i + 5))) ;
for(mapped, evenSink, |out, i|
    if (i % 2 == 0, merge(out, i), out)) ;
for(mapped, oddSink, |out, i|
    if (i % 2 != 0, merge(out, i), out))
```



ROYAL INSTITUTE
OF TECHNOLOGY

Arc Example 3

Scala

```
val mapped = input.map(i: Int => i + 5)
mapped.filter(i: Int => i % 2 == 0).toSink(...)
mapped.filter(i: Int => i % 2 != 0).toSink(...)
```

Arc

```
| source:stream[i32],
evenSink:streamappender[i32],
oddSink:streamappender[i32] |
for(source,
  {evenSink,oddSink},
  |out,i|
    let j = i + 5;
    if (j % 2 == 0,
        {merge(out.$1, j),out.$2},
        {out.$1, merge(out.$2, j)}))
```

Arc Windows

- Windows are supported using higher-order builders (`windower`)
- In addition to *merge-type* and *result-type*, a `windower` also has an *aggregation-type*, which must be another builder
- Use Weld functions to
 - determine window *start* and *end* points,
 - convert the *aggregation-type* into the *result-type*
 - convert the `windower`'s *merge-type* into the aggregation-builder's *merge-type*

Arc Compilation

- Ahead-of-time compilation
- Target long-running jobs (allow costly optimisations)
- Constraint-based type-inference solver
- Compile to a deployment graph IR (and from there to Rust)
- Leave pure Weld expressions to the Weld compiler
- Declarative style allows data *flow* optimisations, such as *operator fusion* and *filter reordering*