

Automerge

Replicated data structures for
peer-to-peer collaboration

Martin Kleppmann, University of Cambridge @martinkl

Joint work with Victor Gomes, Dominic Mulligan, Alastair Beresford, Peter van Hardenberg, Orion Henry, Adam Wiggins, Roshan Choxi, Jeff Peterson, Jim Pick, & more contributors

dataintensive.net

O'REILLY®

Designing Data-Intensive Applications

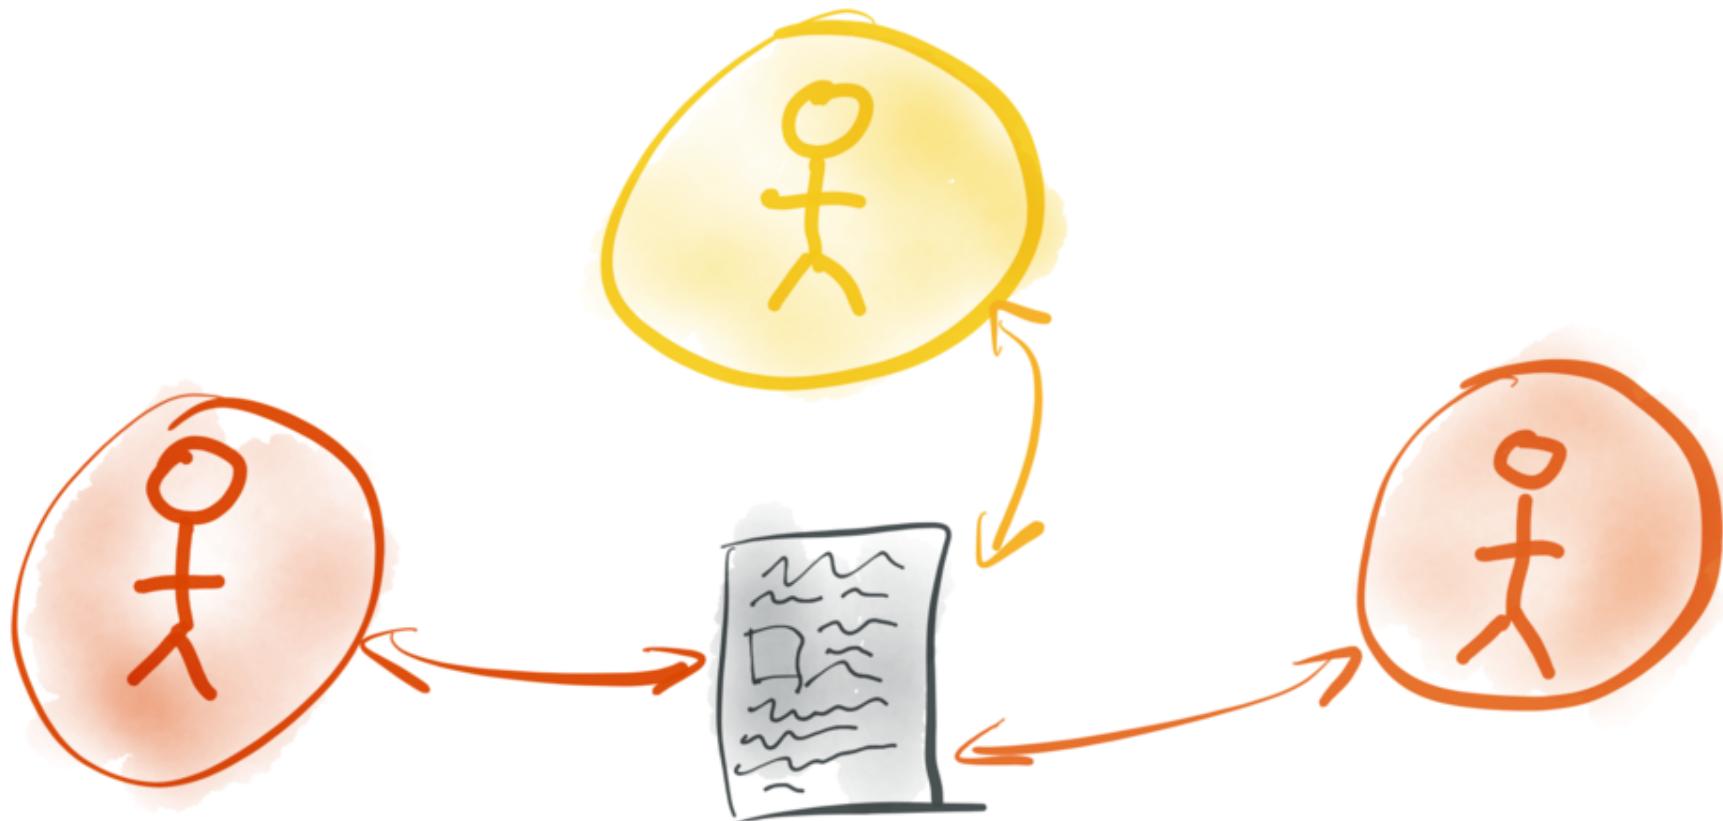
THE BIG IDEAS BEHIND RELIABLE, SCALABLE,
AND MAINTAINABLE SYSTEMS



Martin Kleppmann

@martinkl

COLLABORATIVE APPLICATIONS



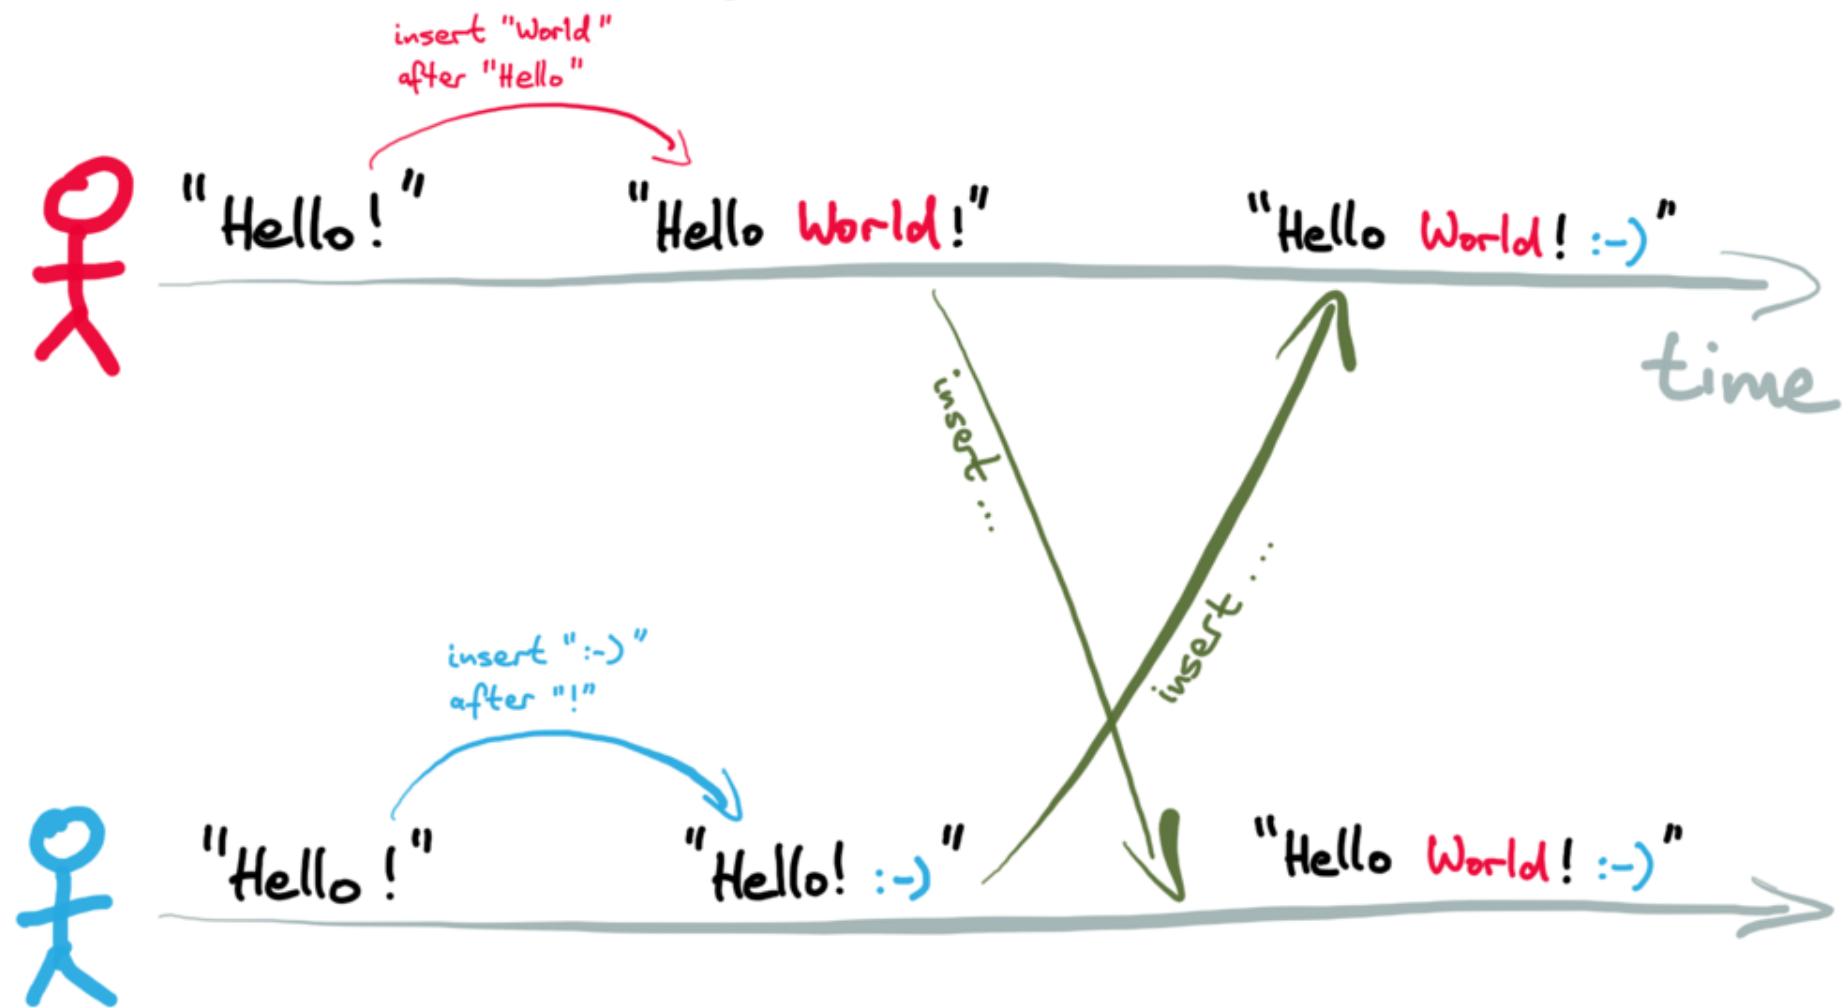
Example : Text editing



Example : Text editing



Example : Text editing



Algorithms for convergence

OPERATIONAL TRANSFORMATION (OT)

- e.g. Google Docs,
MS Office Online

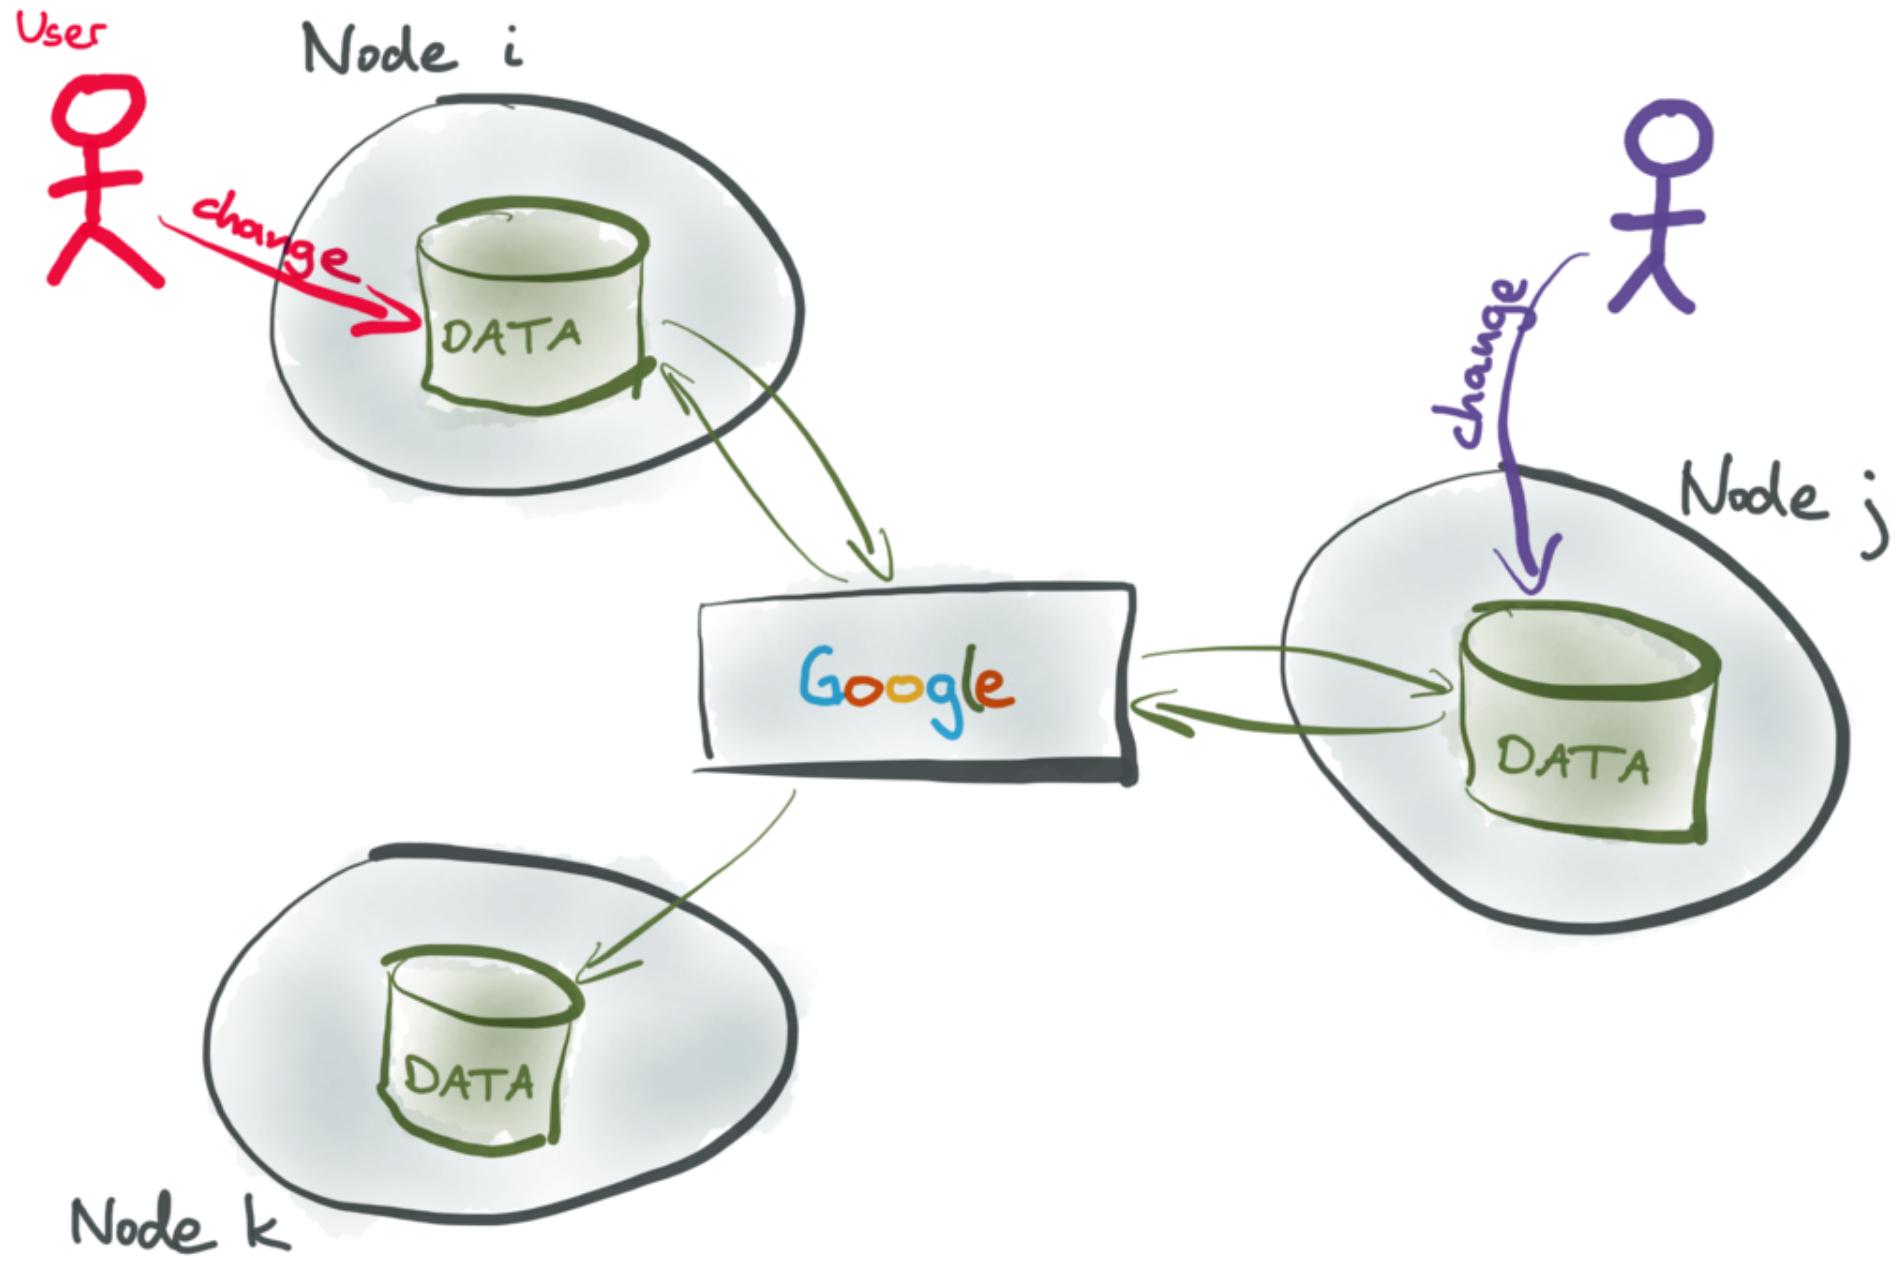
1989–2006

CONFFLICT-FREE REPLICATED DATA TYPES (CRDTs)

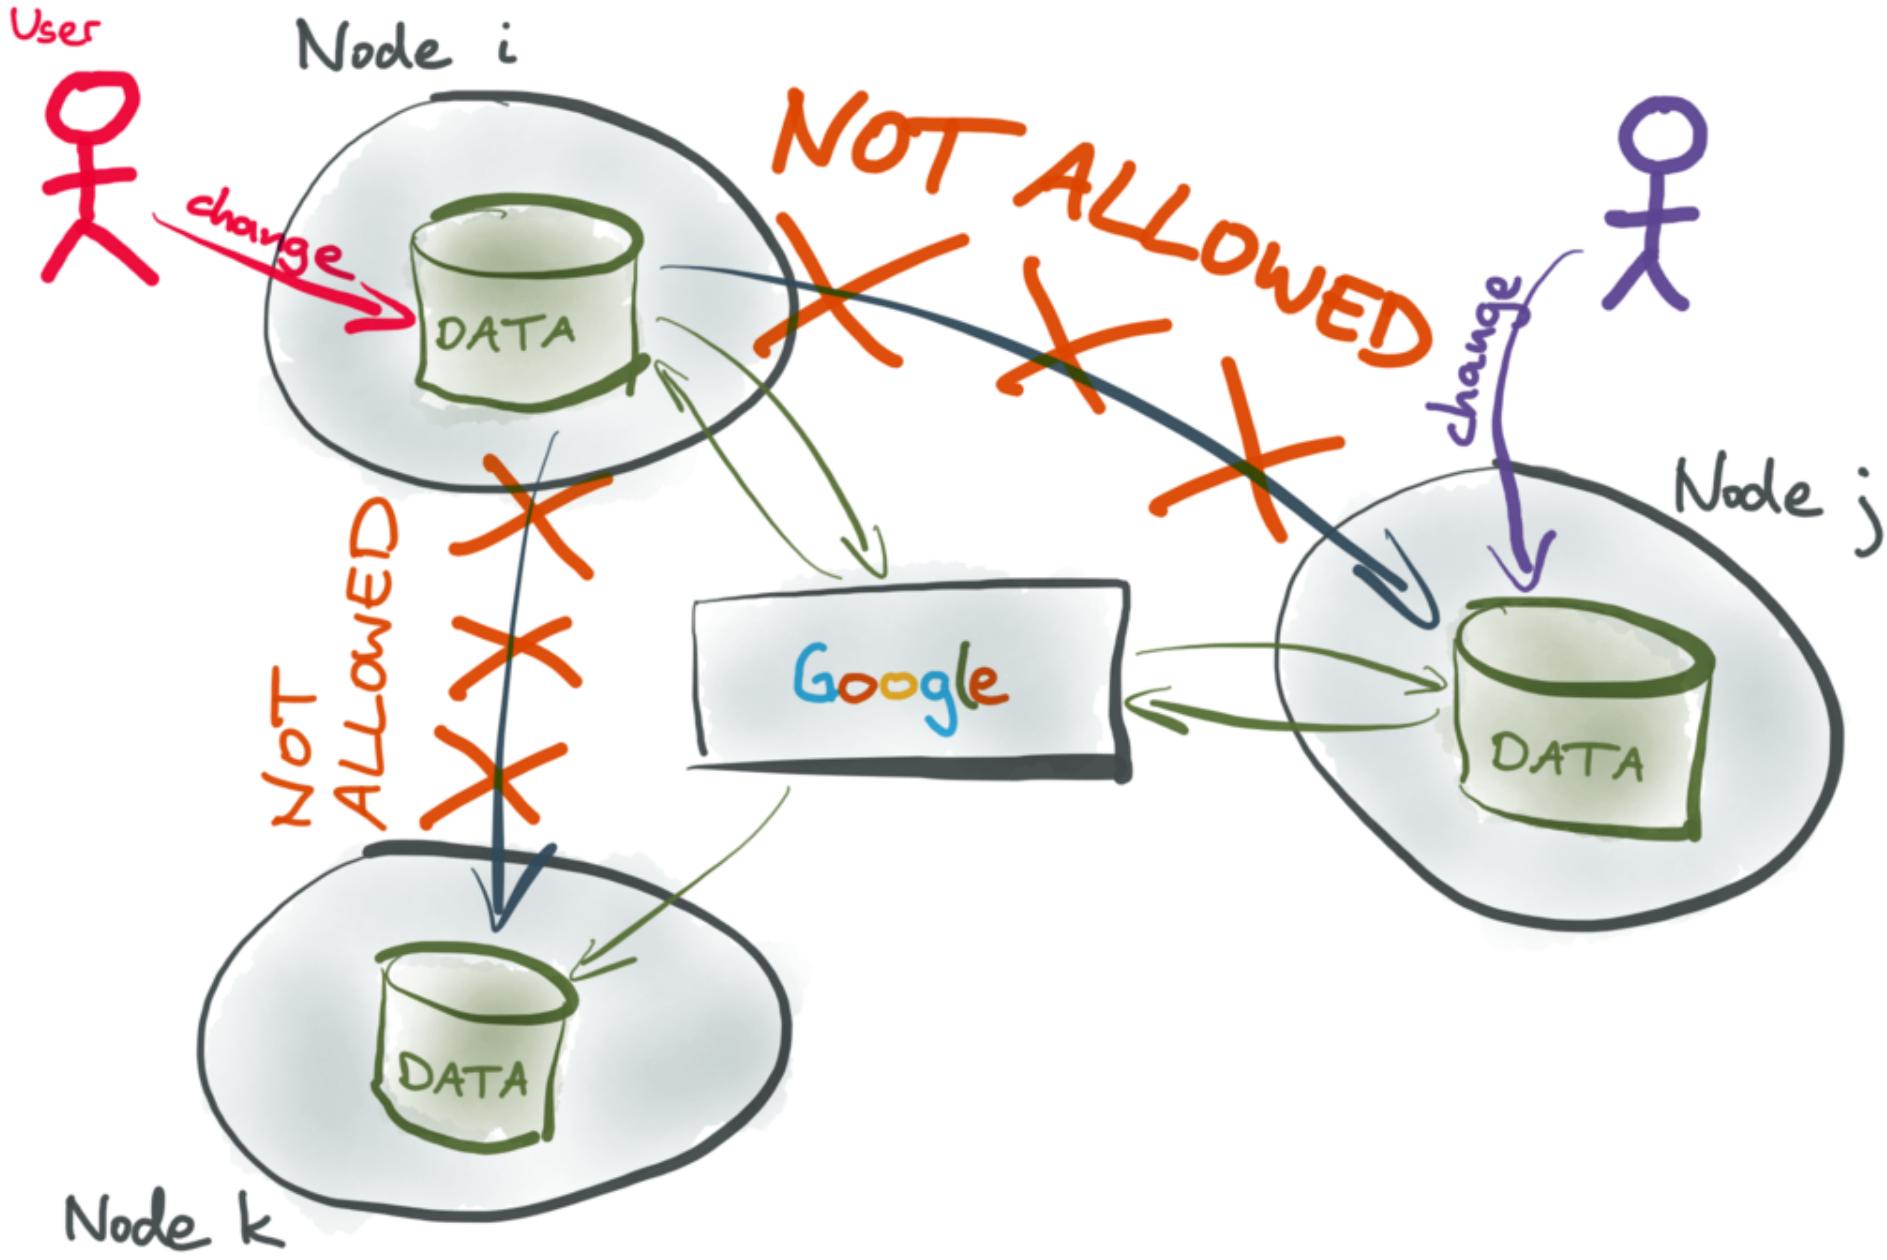
- e.g. Riak, TomTom GPS,
Teletype for Atom, ...

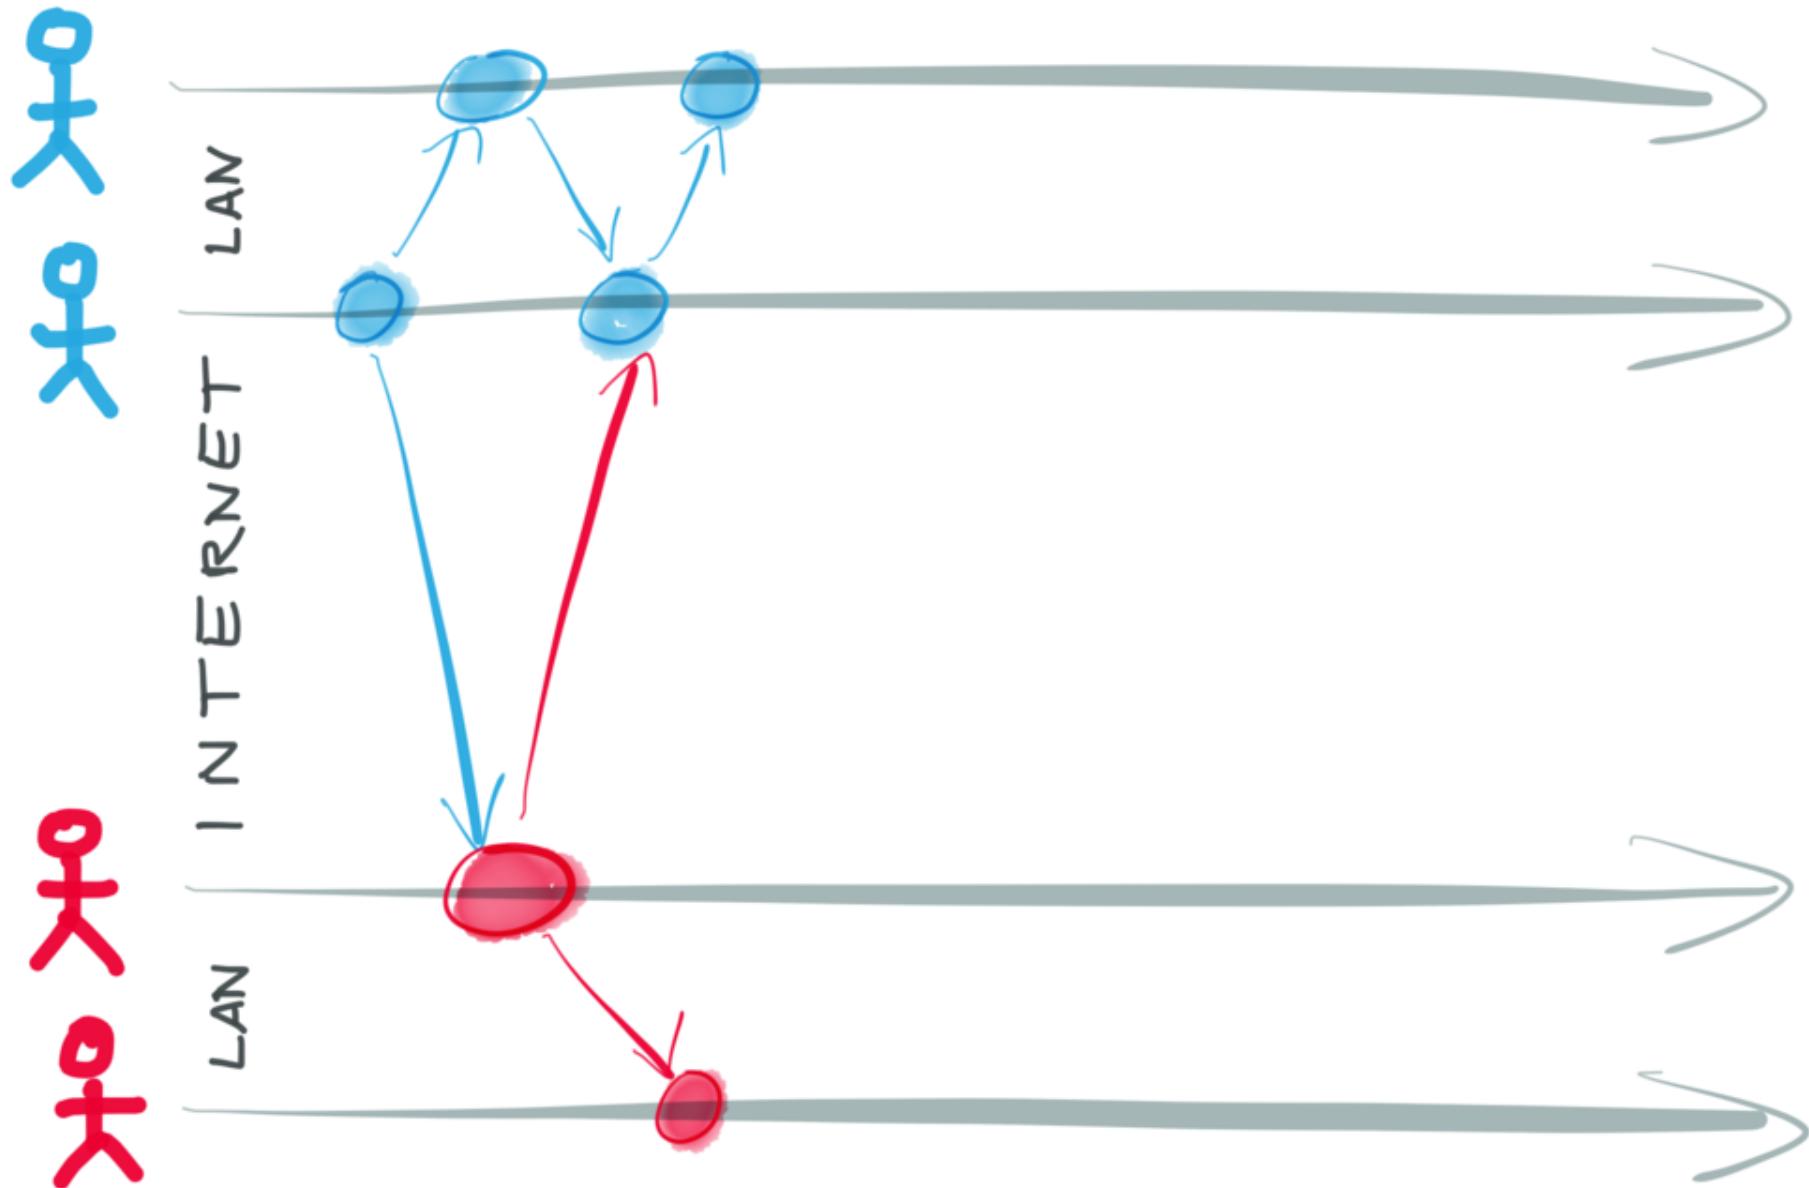
2006 – present

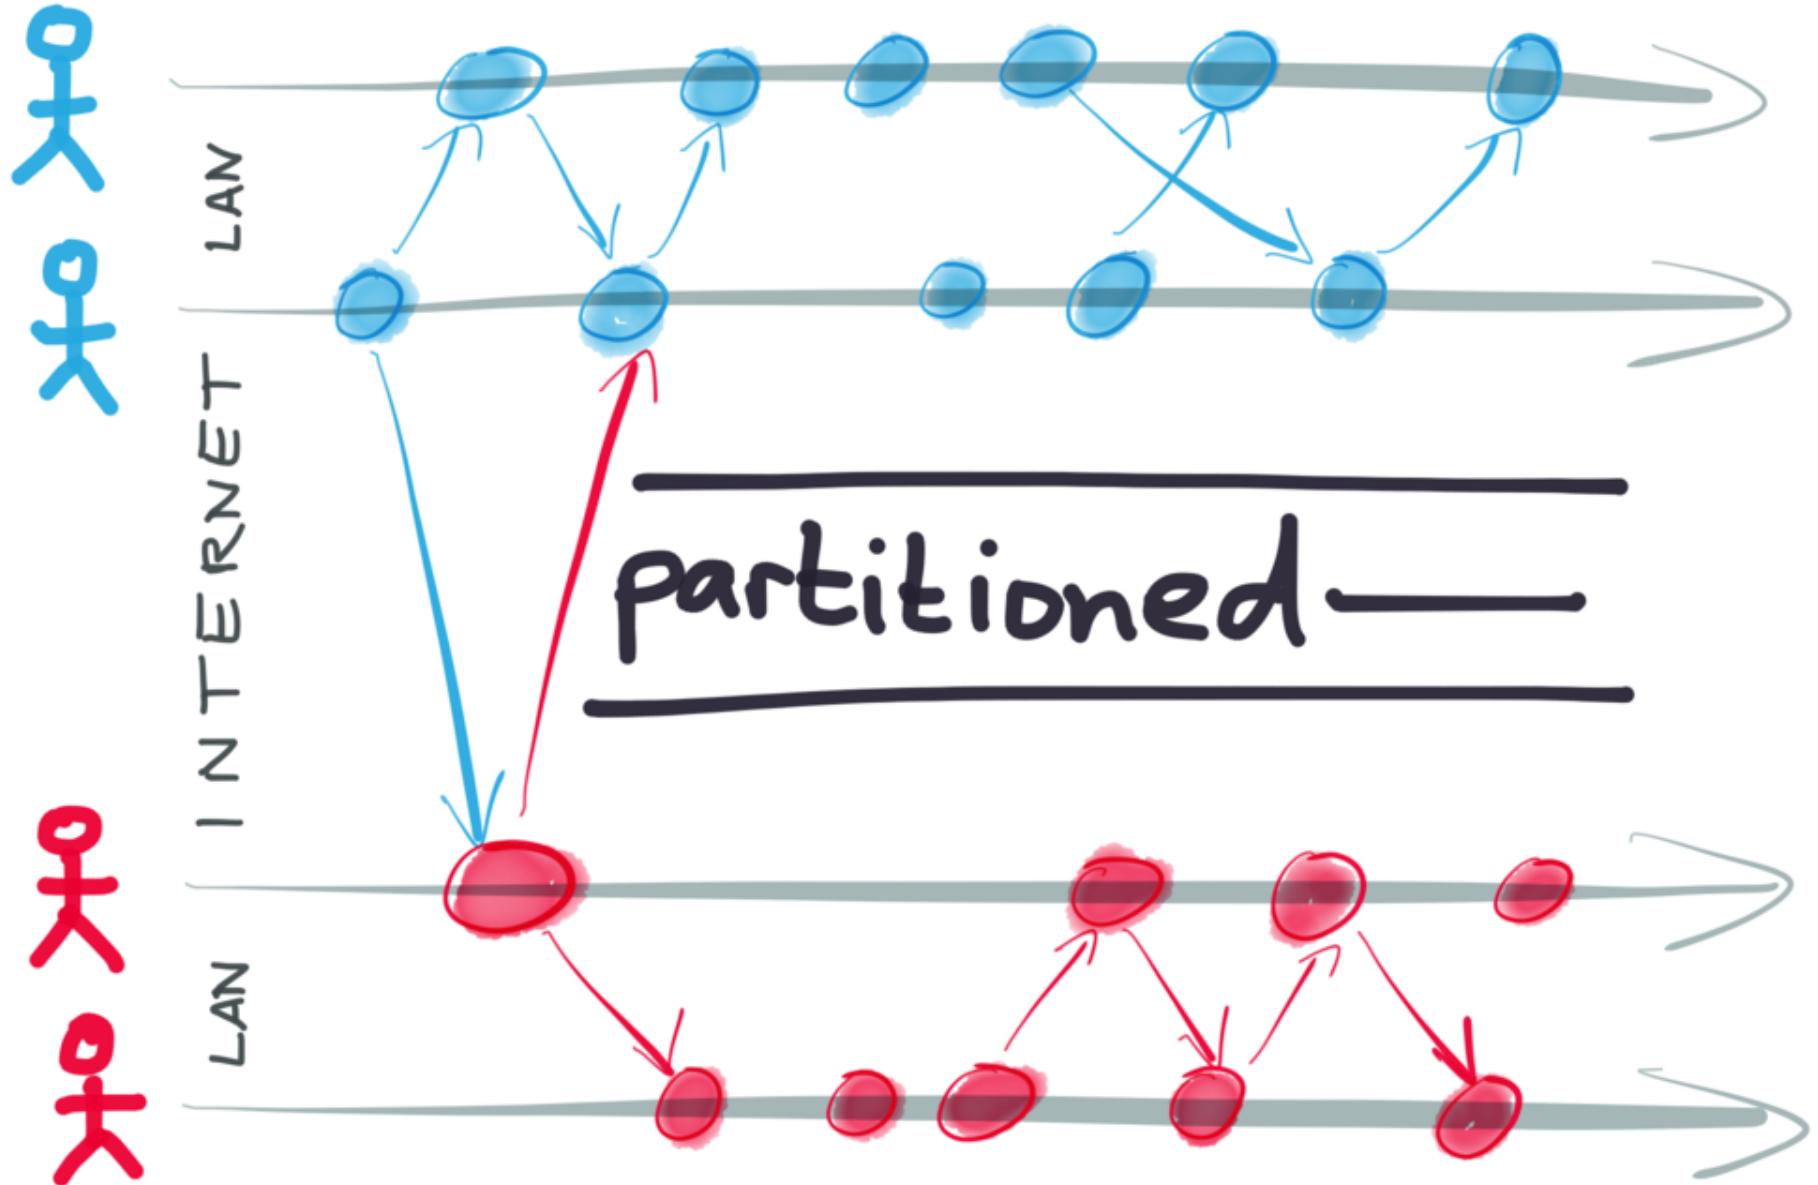
OPERATIONAL TRANSFORMATION IN GOOGLE DOCS.

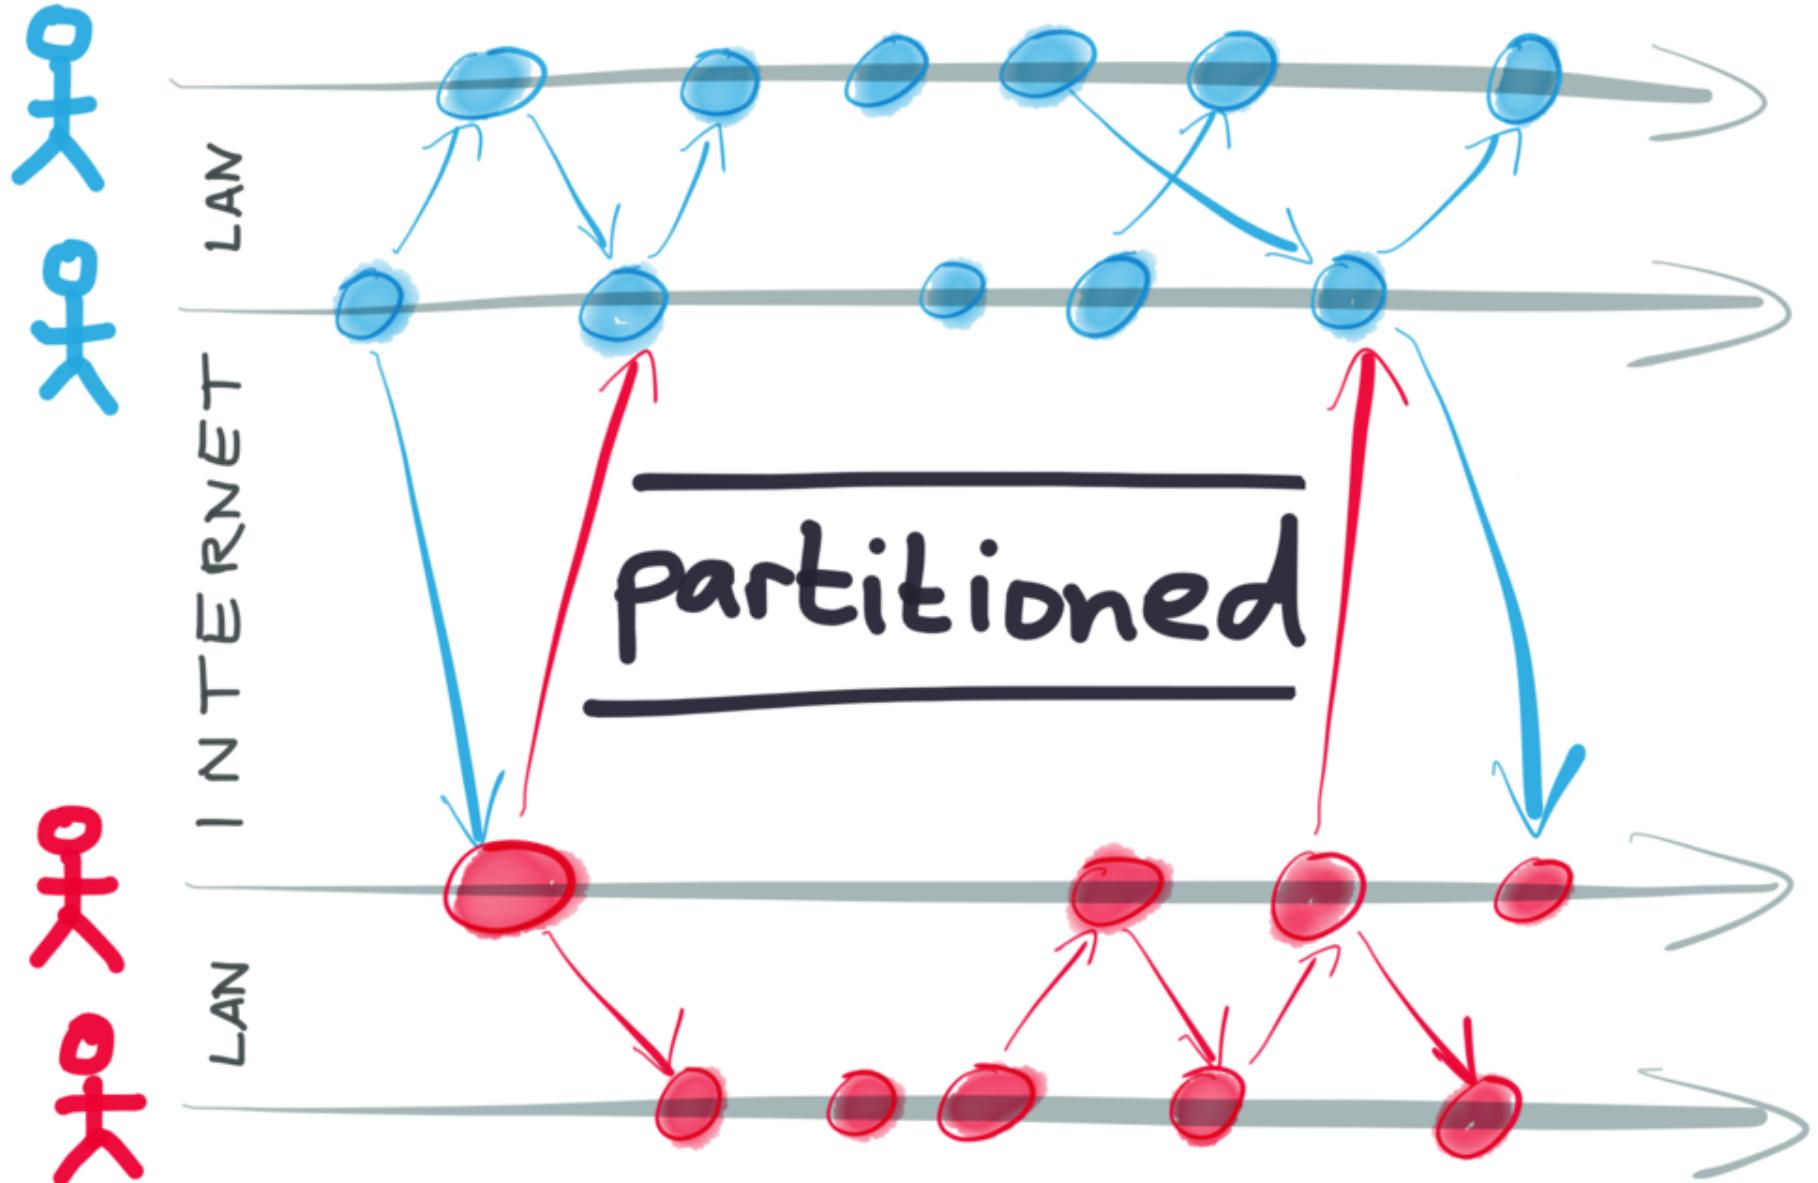


OPERATIONAL TRANSFORMATION IN GOOGLE DOCS.









Algorithms for convergence

OPERATIONAL TRANSFORMATION (OT)

- e.g. Google Docs,
MS Office Online

1989–2006

CONFFLICT-FREE REPLICATED DATA TYPES (CRDTs)

- e.g. Riak, TomTom GPS,
Teletype for Atom, ...

2006 – present

PROVING CRDT's CORRECT

RGA

ORSet

Counter

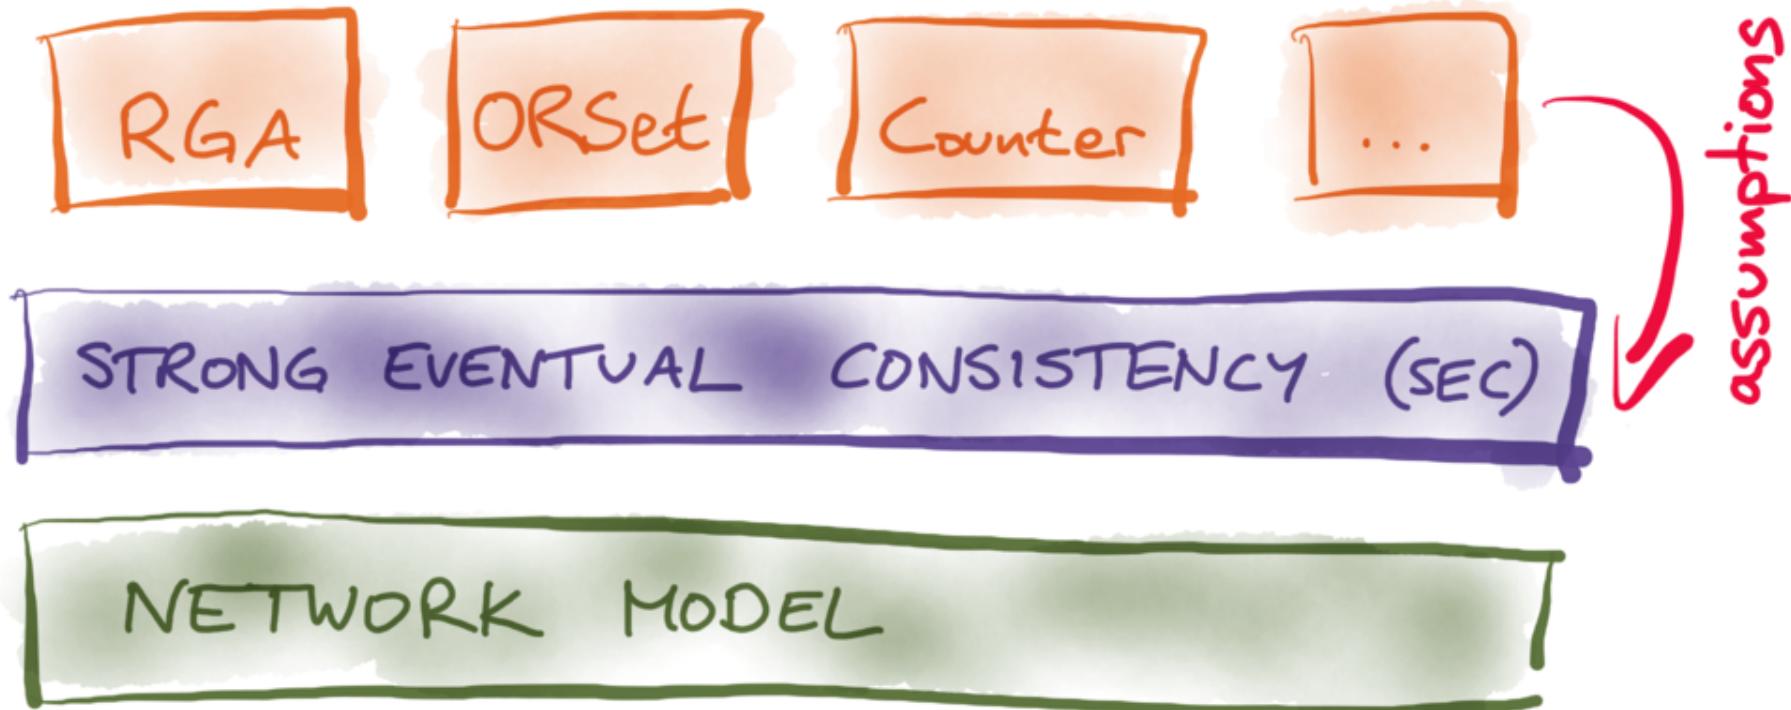
...

STRONG EVENTUAL CONSISTENCY (SEC)

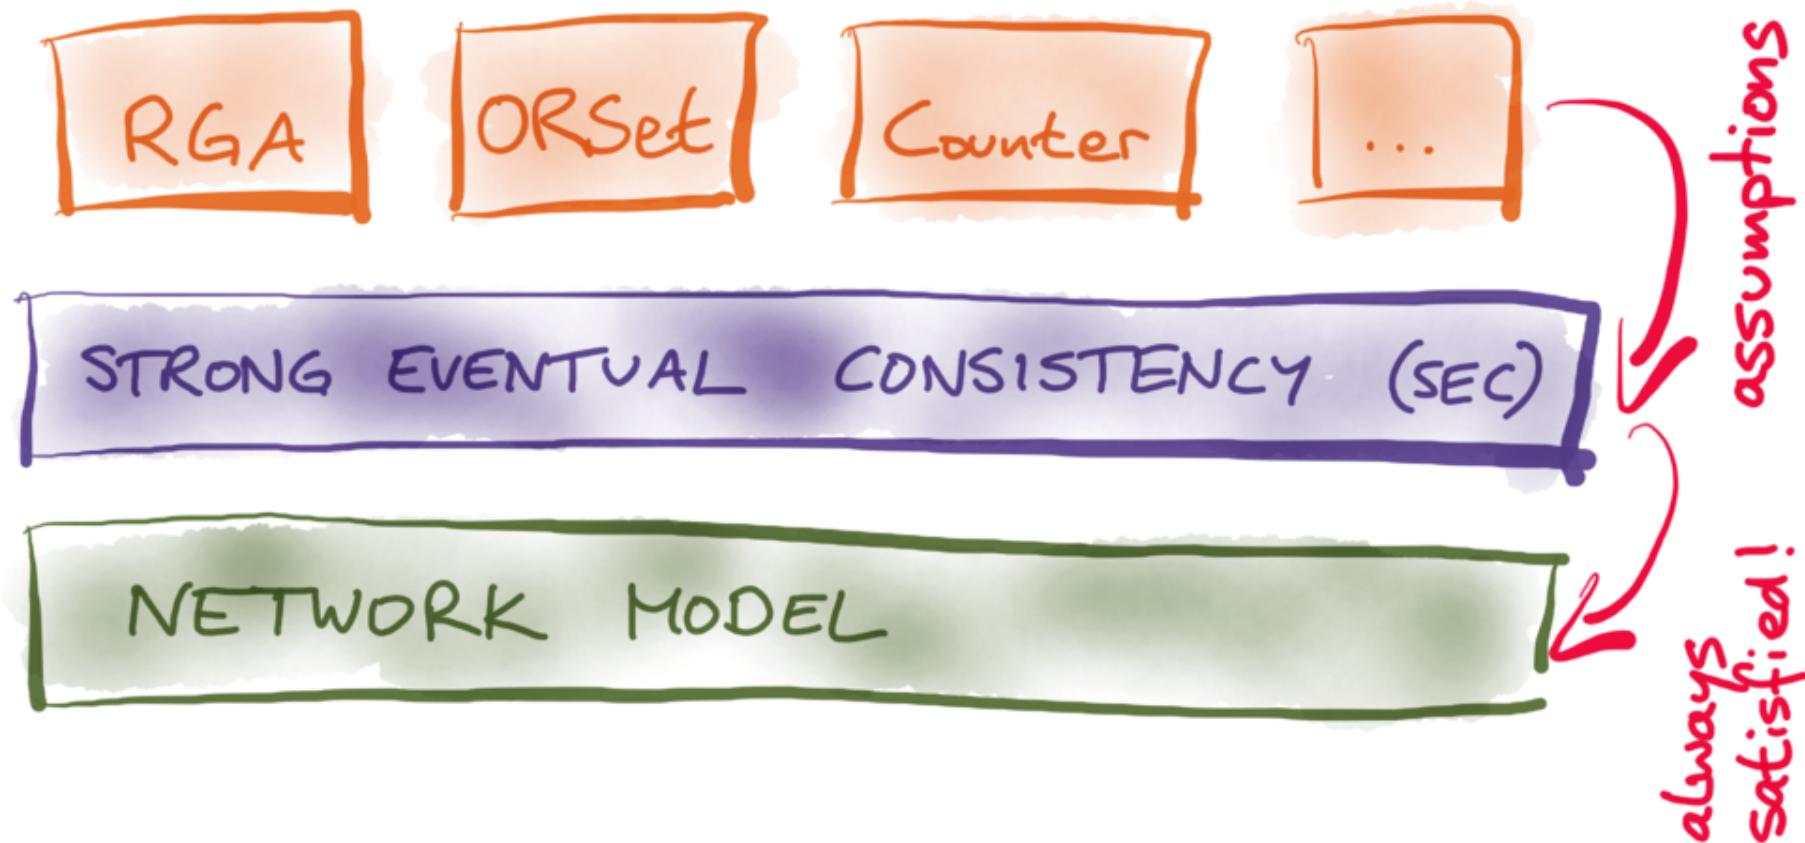
NETWORK MODEL



PROVING CRDT's CORRECT

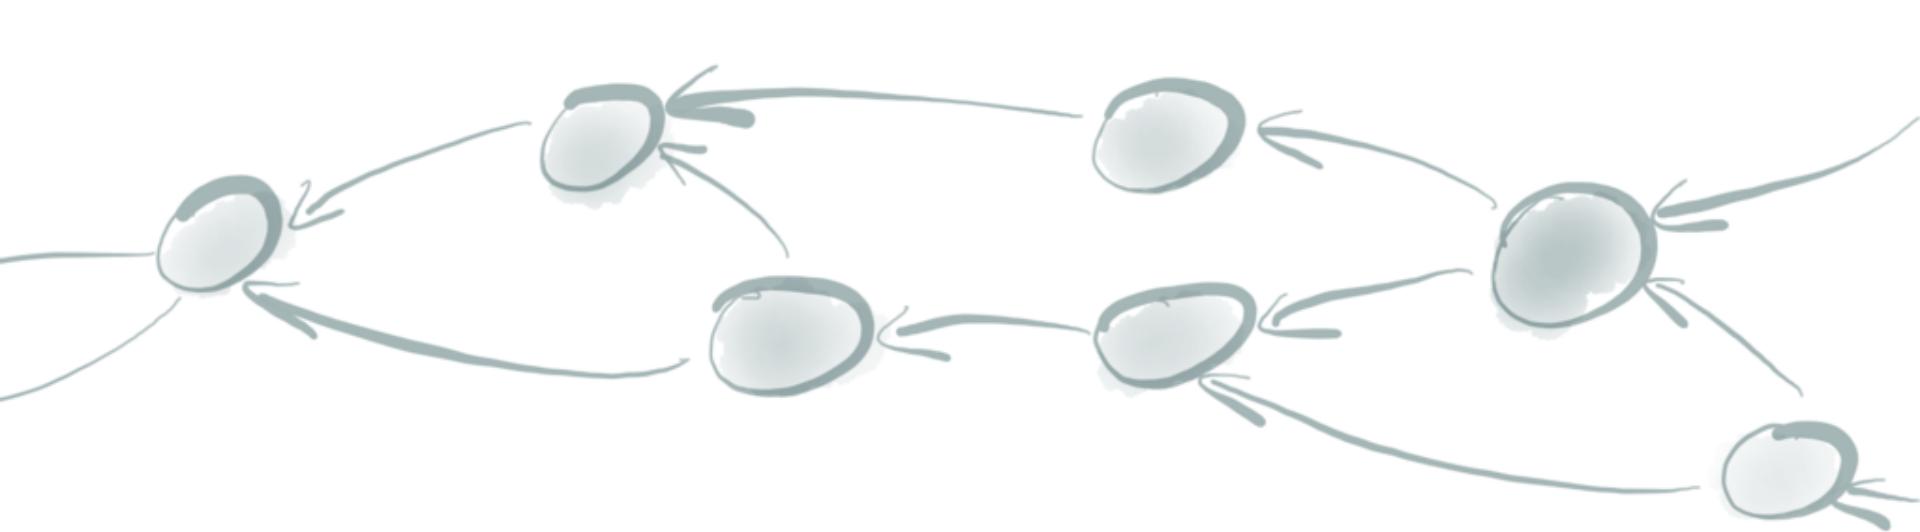


PROVING CRDT's CORRECT



For details, see our paper at <https://doi.org/10.1145/3133933>

Victor B. F. Gomes, Martin Kleppmann, Dominic P. Mulligan, and Alastair R. Beresford:
Verifying Strong Eventual Consistency in Distributed Systems. PACMPL 1(OOPSLA), 2017.



Automerge

<https://github.com/automerge/automerge>

Automerge

(data model / storage)

Apps

Apps

Automerge

(data model / storage)

Apps

Apps

Trellis

Automerge

(data model / storage)

Trellis, a Trello clone based on Automerge: <https://github.com/automerge/trellis>
Joint work with Orion Henry, Peter van Hardenberg, Roshan Choxi, and Adam Wiggins.

The Avengers Initiative

PROSPECTS

RECRUITING

JOINED THE TEAM

Spiderman

1

Iron Man



Doctor Strange

Add a card...

Star Lord

The Hulk

1

Groot

Add a card...

Add a card...

Apps

Apps

Trellis

Pixelpusher

Automerge

(data model / storage)

Pixelpusher, a collaborative pixel art editor: <https://github.com/automerge/pixelpusher>
Created by Javier Valencia, Jeff Peterson, Peter van Hardenberg, and Jim Pick.

PIXEL ART TO CSS

by JVALEN

Star 324



+



+



+



+



25

50

75

100

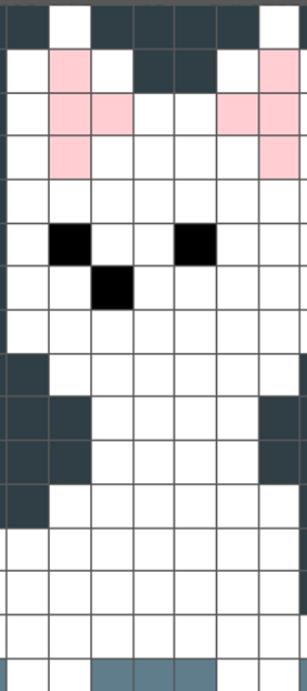
NEW

LOAD

SAVE



css



PREVIEW

RESET

16



16



Pixel Size

10

Duration

1

Apps

Apps

Trellis

Pixelpusher

Automerge

(data model / storage)

MPL

WebRTC

MPL, a WebRTC network layer for Automerge: <https://github.com/automerge/mpl>
Joint work with Orion Henry, Peter van Hardenberg, Roshan Choxi, and Adam Wiggins.

Apps

Apps

Trellis

Pixelpusher

Automerge

(data model / storage)

MPL

Hypermerge

WebRTC

Hypercore/Dat

Hypermerge, a peer-to-peer network layer: <https://github.com/automerge/hypermerge>
Created by Jim Pick, Jeff Peterson, and Peter van Hardenberg.

Apps

Apps

Trellis

Pixelpusher

Automerge

(data model / storage)

MPL

Hypermerge

WebSocket
Client / Server

WebRTC

Hypercore / Dat

APPLICATION DATA (JSON)

```
{ "to-do": [  
    { "title": "buy milk",  
      "done": false },  
    { "title": "water plants",  
      "done": false }  
],  
  "settings": { "alert-sound": "ring", ... }  
}
```

APPLICATION DATA

(JSON)

```
{ "to-do": [  
    { "title": "buy milk",  
      "done": false },  
    { "title": "water plants",  
      "done": false }  
],  
  "settings": {"alert-sound": "ring", ...}  
}
```

ordered list ← "to-do": [] → **list**

APPLICATION DATA

(JSON)

```
{ "to-do": [  
    { "title": "buy milk",  
      "done": false },  
    { "title": "water plants",  
      "done": false }  
],  
  "settings": {"alert-sound": "ring", ...}  
}
```

ordered list

nested maps

EDITING OPERATIONS

```
{ "to-do": [  
    { "title": "buy milk",  
        "done": false },  
    { "title": "water plants",  
        "done": false } ] }  
    "settings": { "alert-sound": "ring", ... }
```

value assignment

true

EDITING OPERATIONS

```
{ "to-do": [ string editing
    { "title": "buy milk",
      "done": false },
    { "title": "water plants",
      "done": false } value assignment
  ],
  "settings": {"alert-sound": "ring", ...}
}
```

The diagram illustrates editing operations on a JSON object. A red cloud highlights the string "milk" in the "title" field of the first item, with the word "soy" written above it, labeled "string editing". Another red cloud highlights the value "false" in the "done" field of the second item, with the word "true" written below it, labeled "value assignment".

EDITING OPERATIONS

```
{"to-do": [
```

{ "title": "buy milk", **string editing**
"done": false }, **list insertion**

{ "title": "water plants",

"done": ~~false~~ **true** } **value assignment**

```
],
```

"settings": {"alert-sound": "ring", ...}

```
}
```

EDITING OPERATIONS

```
{"to-do": [
```

{ "title": "buy milk", **string editing**
"done": false }, **list insertion**

{ "title": "water plants",

"done": ~~false~~ } **value assignment**

},
"settings": {

} **put map key**

"alert-sound": "ring", ... }

"background-image": ...

```
state = Automerge.change(state, "Add todo item",
(doc) => {
    doc.todos.push({
        title: "Buy milk",
        done: false
    })
})

```

state is immutable
(Automerge.change() returns
new object)

state = Automerge.change(state,
"Add todo item",

(doc) => {

doc.todos.push({

title: "Buy milk",

done: false

)

)

state is immutable
(Automerge.change() returns
new object)

"Commit message"
(optional)

state = Automerge.change(state,
(doc) => {

doc.todos.push({

title: "Buy milk",
done: false

)

)

state is immutable
(Automerge.change() returns
new object)

"Commit message"
(optional)

state = Automerge.change(state,

(doc) => {

doc is mutable
only within
this block
(Proxy object)

doc.todos.push({

title: "Buy milk",
done: false

)

)

state is immutable
(Automerge.change() returns
new object)

"Commit message"
(optional)

state = Automerge.change(state,

(doc) => {

doc is mutable
only within
this block
(Proxy object)

doc.todos.push({

title: "Buy milk",
done: false

})

})

Plain old JS
objects and methods

```
doc.todos.push({
```

```
    title: "Buy milk",  
    done: false
```

```
)
```

operation log

```
{action: "makeMap", obj: id1}
```

```
{action: "set", obj: id1, key: "title", value: "Buy milk"}
```

```
{action: "set", obj: id1, key: "done", value: false}
```

```
{action: "ins", obj: todosID, key: prevID, elem: 15}
```

```
{action: "link", obj: todosID, key: elem15, value: id1}
```



Convergence guarantee :

Any two nodes have seen the
same set of operations

(but maybe in a different order!)



They are in the same state

— CONCURRENT CHANGES —

```
{todos: [  
  {title: "Water plants",  
   done: false}  
]}
```

```
{todos: [  
  {title: "Water plants",  
   done: false}  
]}
```

— CONCURRENT CHANGES —

```
{todos: [  
  {title: "Water plants",  
   done: false}  
]}
```

doc.todos.push ({
 title: "Buy milk",
 done: false
})

```
{todos: [  
  {title: "Water plants",  
   done: false}  
]}
```

doc.todos[0].done = true

— CONCURRENT CHANGES —

```
doc.todos.push({  
    title: "Buy milk",  
    done: false  
})
```

```
doc.todos[0].done = true
```

— CONCURRENT CHANGES —

```
doc.todos.push({  
    title: "Buy milk",  
    done: false  
})
```

```
{todos: [  
    {title: "Water plants",  
     done: false},  
    {title: "Buy milk",  
     done: false}  
]}
```

```
doc.todos[0].done = true
```

```
{todos: [  
    {title: "Water plants",  
     done: true}  
]}
```

— CONCURRENT CHANGES —

doc.todos.push({

title: "Buy milk",

done: false

)

{ todos: [

{ title: "Water plants",

done: true },

{ title: "Buy milk",

done: false }

]

doc.todos[0].done = true

network communication

{ todos: [

{ title: "Water plants",

done: true },

{ title: "Buy milk",

done: false }

]

— CONCURRENT CHANGES —

doc.todos[0].title =
"Buy soy milk"

doc.todos[0].title =
"Buy almond milk"

— CONCURRENT CHANGES —

doc.todos[0].title =
"Buy soy milk"

network communication

{ todos : [
 { title : "Buy soy milk",
 done : false,
 -conflicts : { title : {
 1234 : "Buy almond milk"
 }}
 }]}

doc.todos[0].title =
"Buy almond milk"

{ todos : [
 { title : "Buy soy milk",
 done : false,
 -conflicts : { title : {
 1234 : "Buy almond milk"
 }}
 }]}

TIME TRAVEL

Automerge.getHistory(state)

TIME TRAVEL

Automerge.getHistory(state)

```
⇒ [ { change: { message: "Add todo item", ... },  
      snapshot: { todos: [ { title: "Buy milk", ... } ] } },  
    { change: { message: "Mark item as done", ... },  
      snapshot: { todos: [ { title: "Buy milk", ... } ] } },  
    ...  
  ]
```

Resources

- Martin's email: mk428@cl.cam.ac.uk
- Martin on Twitter: [@martinkl](https://twitter.com/martinkl)
- Automerge: <https://github.com/automerge/automerge>
- Trellis: <https://github.com/automerge/trellis>
- Pixelpusher: <https://github.com/automerge/pixelpusher>
- MPL (WebRTC layer): <https://github.com/automerge/mpl>
- Hypermerge: <https://github.com/automerge/hypermerge>
- Dat / Hypercore: <https://datproject.org/>
- Proving CRDTs correct: <https://doi.org/10.1145/3133933>
- JSON CRDT: <http://arxiv.org/abs/1608.03960>
- Martin's book: <http://dataintensive.net/>